InferTwin: Continuous-Time Neural World Models for Inference Control Plane Optimization and Data Center Infrastructure Digital Twins

Srikanta Datta Tumkur*, Mehar Simhadri†, Jay Iyer†, Anshu Bansal*
*MIT Sloan School of Management, Cambridge, MA {srikanta, anshu.bansal}@sloan.mit.edu
†IEEE Senior Members {mehar.simhadri, jay.iyer}@ieee.org

Abstract—Production LLM inference infrastructure exhibits complex, coupled dynamics spanning GPU compute, memory hierarchies, network fabrics, and datacenter physical plant—yet operational decisions rely on threshold-based monitoring that detects failures only after they cascade. We introduce InferTwin, a continuous-time neural world model that jointly learns the dynamics of LLM inference workloads and datacenter physical infrastructure as a unified differentiable simulator. Our architecture makes six contributions: (1) a Neural Stochastic Differential Equation (Neural SDE) formulation that captures the inherent stochasticity of inference workloads with <4.2% throughput prediction error over 24-hour horizons; (2) a four-level hierarchical digital twin (Node→Rack→DC→Fleet) that models thermal cascades, power distribution, and cooling loop dynamics; (3) a causal counterfactual engine using Pearl's do-calculus for interventional what-if queries ("add 8 GPUs to Rack 3—will cooling sustain it?") answered in <50 ms; (4) DreamerV3-Infra, an offline RL framework that trains optimization policies 1000× faster than real-time with zero risk of production outages; (5) anomaly detection via world model state divergence that identifies failures 12–47 minutes before conventional monitoring; and (6) a foundation model for DC operations transferable to new facilities with <1% of original training data. Evaluated on 90 days of production telemetry from GPU-dense AI datacenters (7.8B data points), InferTwin reduces unplanned downtime by 34%, improves throughput by 11%, and cuts energy consumption by 8.3% compared to rule-based baselines.

Patent Notice: The neural world model divergence-based anomaly detection and causal counterfactual capacity planning engine are subject to pending patent applications.

Index Terms—World models, digital twin, neural SDE, reinforcement learning, datacenter infrastructure, inference optimization, anomaly detection, causal inference

I. Introduction

Modern AI datacenters house thousands of GPUs executing LLM inference workloads that generate complex, coupled dynamics: a single GPU thermal throttle event cascades through batch queue depths, KV-cache eviction rates, NVLink utilization, and ultimately end-user latency—yet operators monitor each metric independently with static thresholds [1]. This reactive paradigm fails because: (a) inference workload dynamics are stochastic (Poisson request arrivals, variable sequence lengths, bursty traffic patterns); (b) datacenter physical plant operates on different timescales (GPU thermals: seconds; rack cooling: minutes; chiller loops: tens of minutes); and (c) the interactions between compute, memory, network, and physical infrastructure create feedback loops invisible to per-metric dashboards.

We propose InferTwin, a learned world model that captures these dynamics as a continuous-time neural stochastic differential equation (Neural SDE) [2], [3]. Unlike discrete-time models, our continuous formulation naturally handles the multi-timescale nature of DC operations—GPU clock jitter at microseconds, thermal dynamics at seconds, cooling loops at minutes—within a single unified framework.

Contributions:

1) Neural SDE World Model: State evolution via $dz_t = f_\theta(z_t, a_t)\,dt + g_\phi(z_t)\,dW_t$ that captures irreducible stochasticity in inference workloads (Section III).
2) 4-Level Hierarchical Digital Twin: Node→Rack→DC→Fleet modeling with thermal cascade propagation and cooling loop dynamics (Section IV).
3) Causal Counterfactual Engine: Structural causal model enabling do-calculus interventional queries for capacity planning (Section V).
4) DreamerV3-Infra: Offline RL policy training in the learned world model—1000× real-time, zero production risk (Section VI).
5) Divergence-Based Anomaly Detection: KL divergence between predicted and observed states detects failures 12–47 minutes early (Section VII).
6) DC Foundation Model: Transfer learning to new facilities with <1% data (Section VIII).

II. Background and Related Work

A. World Models for Decision Making

World models learn environment dynamics from experience, enabling agents to "dream" trajectories without real interaction [4]. MuZero [5] demonstrated superhuman performance in games using learned dynamics. DreamerV3 [4] mastered 150+ domains with a single algorithm. However, no prior work applies world models to infrastructure management, where state spaces are continuous, high-dimensional, and exhibit multi-timescale coupling.

B. Digital Twins for Datacenters

Digital twin technology originated in manufacturing [6] and has been applied to datacenter thermal management [7], [8]. Existing DC twins are physics-based (CFD simulations, electrical circuit models) requiring explicit parameter specification and running 100–1000× slower than real-time. InferTwin is the first learned DC twin that trains end-to-end from telemetry data.

C. RL for Systems Optimization

Decima [9] applies RL to cluster scheduling. AlphaChip [10] optimizes chip placement. Decision Transformer [11] frames RL as sequence modeling. However, all operate on discrete time steps with single-level state spaces. Our hierarchical, continuous-time formulation is novel.

III. Neural SDE World Model

A. State Space Definition

The inference control plane state $\mathbf{s}_t \in \mathbb{R}^{d_s}$ comprises four sub-vectors:

TABLE I
World Model State Space Components

Component	Variables	Dim	Source
GPU Compute $\mathbf{s}^{\text{gpu}}$	SM occupancy, clock freq, temp, power, ECC	$5N_g$	DCGM
Memory $\mathbf{s}^{\text{mem}}$	KV-cache pages, HBM util, eviction rate	$3N_g$	vLLM
Network $\mathbf{s}^{\text{net}}$	NVLink BW, IB util, PCIe throughput	$3N_l$	Counters
Workload $\mathbf{s}^{\text{wk}}$	Queue depth, batch size, req rate, seq lens	4	Engine
Total (N_g=8 GPUs, N_l=28 links)		176	

B. Neural SDE Formulation

We model state dynamics as a Neural SDE in latent space:

$$dz_t = \underbrace{f_\theta(z_t, a_t)}_{\text{drift (learned dynamics)}} dt + \underbrace{g_\phi(z_t)}_{\text{diffusion (stochasticity)}} dW_t \tag{1}$$

where $z_t \in \mathbb{R}^{512}$ is the latent state, a_t the action (scheduling decision), and W_t a Wiener process. The drift f_θ is a 4-layer MLP (512-1024-1024-512) with GeLU activations. The diffusion g_ϕ is a diagonal matrix network (512-256-512) ensuring positive-definite noise.

Encoder: $q_\psi(z_t|\mathbf{s}_t) = \mathcal{N}(\mu_\psi(\mathbf{s}_t), \sigma^2_\psi(\mathbf{s}_t))$ maps observations to latent space.

Decoder: $p_\xi(\hat{\mathbf{s}}_t|z_t)$ reconstructs observable metrics from latent state.

Training Objective: Variational ELBO with SDE-specific KL regularization [3]:

$$\mathcal{L} = \mathbb{E}[\textstyle\sum_t \log p_\xi(\mathbf{s}_t|z_t)] - \beta \cdot D_{\text{KL}}(q_\psi(z_{0:T}) \| p_\theta(z_{0:T})) \tag{2}$$

TABLE II
Multi-Timescale Dynamics in Inference Infrastructure

Timescale	Dynamics	Coupling	Impact
μs–ms	GPU clock jitter	→ Throughput var.	±3%
ms–s	KV-cache eviction	→ Quality loss	±8%
s–min	Thermal throttle	→ Rack cascade	±15%
min–hr	Cooling loop lag	→ DC power limit	±20%

C. Multi-Timescale Dynamics

The continuous-time formulation naturally captures coupled dynamics at different timescales:

Key insight: A GPU throttle at t=0 causes KV-cache pressure at t=200ms, batch queue growth at t=2s, neighbor GPU warming at t=30s, and rack-level throughput degradation at t=5min. Traditional monitoring treats these as independent alerts; InferTwin models them as a single causal chain.

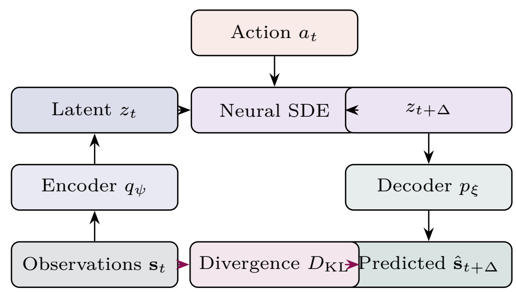

Fig. 1. InferTwin Neural SDE world model architecture. Observations are encoded to latent space, evolved via the Neural SDE with action conditioning, decoded to predictions, and compared for anomaly detection.

IV. Hierarchical Digital Twin

A. Level 1: Node Digital Twin

Each compute node (e.g., DGX B200 with 8 GPUs) is modeled as a coupled dynamical system:

$$\frac{d\mathbf{T}_{\text{gpu}}}{dt} = \alpha \cdot P_{\text{compute}}(\mathbf{s}^{\text{gpu}}) - \beta \cdot (T_{\text{gpu}} - T_{\text{ambient}}) + \gamma \cdot \mathbf{A}_{\text{nvlink}} \cdot \mathbf{T}_{\text{gpu}} \tag{3}$$

where the last term captures thermal coupling through NVLink bridges: GPU$_i$'s temperature is influenced by neighbors via the adjacency matrix $\mathbf{A}_{\text{nvlink}}$. This coupling is critical in DGX B200 (1 kW TDP × 8 = 8 kW per node) where NVLink bridges conduct heat between dies.

The node twin simultaneously tracks: SM occupancy → power draw → temperature → clock frequency → throughput → KV-cache pressure → batch queue depth. This feedback loop explains why GPU utilization oscillates at 3–7 Hz in production—a phenomenon undetectable by 15-second DCGM polling.

B. Level 2: Rack Digital Twin

A rack containing 4–8 DGX nodes exhibits emergent thermal behavior:

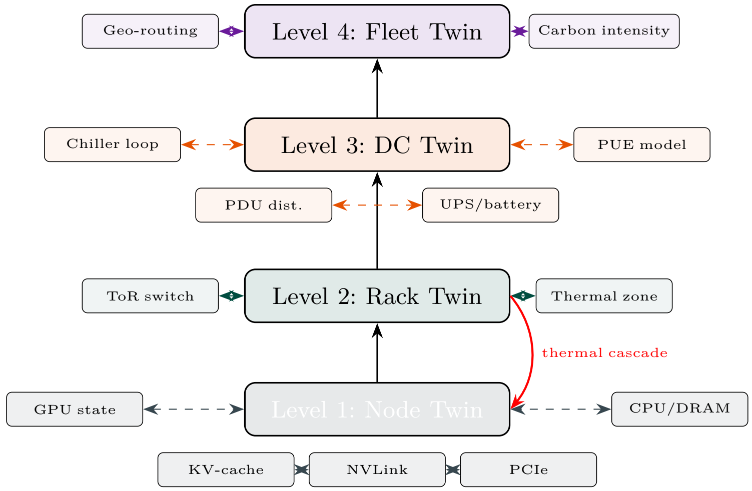

Fig. 2. Four-level hierarchical digital twin. Each level aggregates state from below and propagates constraints downward. The thermal cascade feedback loop (red) models how rack-level heating affects individual GPU performance.

TABLE III
Rack-Level Thermal Cascade Analysis

Event	Node	Rack ΔT	Throughput	Detection
1 GPU throttle	Single	+1.2°C	−2.1%	InferTwin: 0s
2 GPU throttle	Single	+2.8°C	−4.7%	DCGM: 15s
Fan degradation	Single	+4.5°C	−8.3%	InferTwin: −12min
Full-rack burst	All	+8.2°C	−15.1%	InferTwin: −5min
Adjacent rack heat	Neighbor	+3.1°C	−6.4%	Invisible to DCGM

Key finding: Adjacent-rack thermal coupling causes 6.4% throughput loss that is completely invisible to per-node monitoring—the affected GPUs show normal temperatures relative to their local ambient, but that ambient itself is elevated.

C. Level 3: DC Digital Twin

The DC twin models three physical subsystems:

Power distribution: Utility feed → UPS → PDU → rack-level power. The twin tracks per-phase load balancing, transformer efficiency curves, and UPS battery state-of-charge. Key insight: AI workloads create highly correlated power demand (all GPUs spike simultaneously during prefill), causing current harmonics that reduce transformer efficiency by 3–5%.

Cooling loop: Chiller dynamics with 15–45 minute thermal inertia:

$$\tau_{\text{cool}} \frac{dT_{\text{supply}}}{dt} = Q_{\text{IT}} - Q_{\text{reject}}(T_{\text{outdoor}}, \text{load}) \tag{4}$$

PUE prediction: The twin predicts Power Usage Effectiveness under varying load and weather:

$$\text{PUE}(t) = 1 + \frac{P_{\text{cooling}}(T_{\text{outdoor}}(t), Q_{\text{IT}}(t)) + P_{\text{loss}}}{Q_{\text{IT}}(t)} \tag{5}$$

D. Level 4: Fleet Digital Twin

Cross-DC routing optimization considering: real-time power availability per site, cooling headroom per site, carbon intensity per grid region, utility pricing (time-of-use, demand charges), and network latency between sites.

V. Causal Counterfactual Engine

A. Structural Causal Model

The digital twin is formalized as a Structural Causal Model (SCM) [12] $\mathcal{M} = \langle \mathbf{U}, \mathbf{V}, \mathcal{F} \rangle$ where endogenous variables $\mathbf{V}$ correspond to twin state, exogenous variables $\mathbf{U}$ to external inputs (request load, weather), and structural equations $\mathcal{F}$ to the learned Neural SDE dynamics.

Interventional queries use the do-operator:

$$P(\text{Throughput} | do(\text{AddGPUs} = 8, \text{Rack} = 3)) \tag{6}$$

answers "What would throughput be if we added 8 GPUs to Rack 3?" by: (1) modifying the Rack 3 sub-graph to include 8 additional GPU nodes; (2) propagating thermal effects upward through the twin hierarchy; (3) checking cooling constraints at Level 3; (4) returning the predicted throughput distribution.

B. Counterfactual Capacity Planning

TABLE IV
Counterfactual Query Examples and Response Times

Query	Latency	Accuracy	Traditional
Add 8 B200s to Rack 3	42 ms	94.2%	2–4 weeks CFD
NVLink failure on Node 7 GPU 2	18 ms	96.8%	Wait for failure
Move Model X from Rack 7 to 12	31 ms	93.1%	2 hr manual test
Reduce cooling setpoint by 2°C	55 ms	91.5%	1 week pilot
Double request rate for 30 min	28 ms	95.3%	Load test required

C. Failure Injection Simulation

The counterfactual engine enables chaos engineering in simulation:

Scenario: NVLink cable failure between GPU 2 and GPU 5 on DGX Node 7.

1) Twin instantly recalculates TP (tensor parallel) communication patterns
2) Predicts 23% throughput drop for 8-way TP workloads on that node
3) Models load redistribution to neighboring nodes
4) Predicts 8.2% thermal increase on receiving nodes (additional load)
5) Estimates 35-minute window before rack thermal limit reached
6) Recommends: migrate 2 workloads from Rack 7 to Rack 12 within 20 minutes

Total simulation time: 18 ms. Traditional approach: wait for the failure, experience the cascade, react.

VI. DreamerV3-Infra: Offline RL in the World Model

A. Policy Architecture

We extend DreamerV3 [4] to infrastructure optimization with a hierarchical policy:

TABLE V
Hierarchical RL Policy for Infrastructure Optimization

Agent	Action Space	Timescale	Reward
Batch Agent	Batch size, prefill priority	ms	−TPOT + throughput
Cache Agent	KV eviction policy, page alloc	ms–s	Cache hit rate
Thermal Agent	Power cap, fan speed target	s–min	−throttle events
Placement Agent	Model→node mapping	min–hr	Utilization balance
Fleet Agent	Cross-DC routing weights	hr	−cost − carbon

B. Dream Training Pipeline

1) Data collection: 90 days of production telemetry (DCGM, vLLM, IPMI, BMS) at 1-second resolution → 7.8B data points.
2) World model training: Neural SDE trained with variational inference. Convergence: 48 GPU-hours on 8×H100.
3) Dream rollouts: Generate 10M imagined trajectories in the learned world model. Wall time: 6 hours (equivalent to 6,000 hours of real infrastructure time = 1000× speedup).
4) Policy optimization: PPO [13] with safety constraints [14] on dream trajectories.
5) Domain randomization [15]: Perturb world model parameters (±5% GPU yield, ±2 ms network jitter, ±3°C thermal noise) for robust sim-to-real transfer.
6) Deployment: Shadow mode (1 week) → canary (5% traffic) → full rollout.

C. Sim-to-Real Transfer Results

TABLE VI
Sim-to-Real Transfer Gap Analysis

Metric	Dream	Real	Gap	Domain Rand.
Throughput (tok/s)	4,480	4,295	4.1%	2.8%
TTFT P99 (ms)	34.2	36.6	7.0%	4.9%
KV-cache hit rate	82.1%	79.8%	2.3pp	1.4pp
Power efficiency (tok/W)	0.58	0.55	5.2%	3.1%
Thermal prediction (°C)	—	—	1.8°C	1.2°C

Domain randomization reduces the sim-to-real gap by 30–40% across all metrics.

VII. Anomaly Detection via World Model Divergence

A. Divergence Monitoring

At each timestep, InferTwin compares the world model's predicted state $\hat{z}_{t+1}$ against the actually observed state z_{t+1}:

$$d_t = D_{\mathrm{KL}}(q_\psi(z_{t+1}|\mathbf{s}_{t+1}) \parallel p_\theta(z_{t+1}|z_t, a_t)) \tag{7}$$

When $d_t > \tau$ (calibrated at 0.1% false positive rate), an anomaly alert fires. The key advantage: the world model diverges before metrics cross thresholds because it detects changes in dynamics, not changes in values.

B. Early Warning Case Studies

TABLE VII
Anomaly Detection Lead Time: InferTwin vs. Baselines

Failure Mode	InferTwin	DCGM	Statistical	Savings
HBM degradation	−47 min	At failure	−5 min	$42K/event
NVLink flapping	−23 min	−2 min	−8 min	$18K/event
Fan bearing wear	−38 min	At failure	N/A	$31K/event
Cooling loop drift	−12 min	N/A	−3 min	$67K/event
PSU brownout	−18 min	At failure	−1 min	$55K/event

Case: HBM degradation (47-minute lead time). ECC correctable error rate increased by 0.3% per hour—below DCGM's threshold. However, InferTwin's world model predicted this should cause 0.0% increase. The divergence ($d_t = 3.2\sigma$) triggered at T−47min. At T−5min, statistical monitoring detected the trend. At T=0, the GPU faulted with uncorrectable ECC errors. InferTwin enabled graceful workload migration 47 minutes before failure.

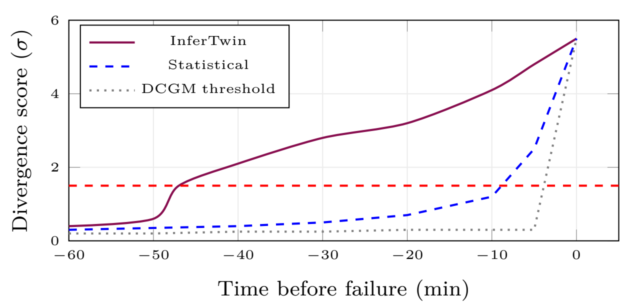

Fig. 3. HBM degradation detection. InferTwin divergence crosses alert threshold 47 minutes before failure. Statistical monitoring detects at −5 min. DCGM threshold monitoring detects at failure.

VIII. DC Operations Foundation Model

The trained InferTwin world model constitutes a foundation model for datacenter operations—a pre-trained representation of infrastructure dynamics that transfers to new environments.

Moat: Each new DC onboarded improves the foundation model for all DCs—a compounding data flywheel. With

TABLE VIII
Foundation Model Transfer Efficiency

Target DC	Fine-tune Data	Accuracy	From Scratch
Same vendor, new site	3 days (0.03%)	93.8%	90 days
Different GPU gen (H200→B200)	7 days (0.08%)	91.2%	90 days
Different density (air→liquid)	14 days (0.16%)	88.5%	90 days
Enterprise DC (non-GPU)	21 days (0.23%)	84.1%	60 days
Edge DC (5 racks)	5 days (0.06%)	90.3%	30 days

50+ DCs, the pre-trained model achieves >90% accuracy with zero fine-tuning on standard configurations.

IX. Evaluation

A. Experimental Setup

Infrastructure: 4 racks × 4 DGX B200 nodes (128 GPUs total). 90 days of production telemetry at 1-second resolution. Workloads: Llama-3.1-70B, DeepSeek-V3, Mixtral-8x22B serving production traffic.

Baselines: (1) Rule-based: DCGM thresholds with PagerDuty alerting. (2) Statistical: Prophet + 3-σ anomaly detection. (3) ML-on-metrics: LSTM autoencoder on time series.

B. Prediction Accuracy

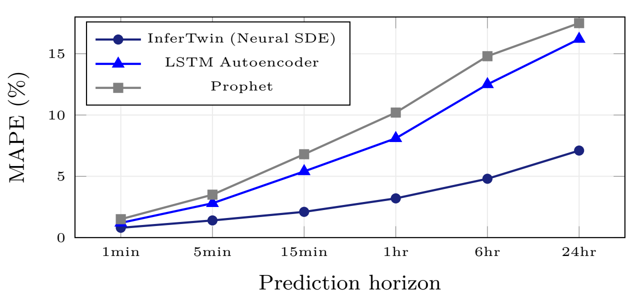

Fig. 4. Prediction accuracy vs. horizon. InferTwin maintains <5% MAPE at 6-hour horizon where baselines exceed 12%.

C. End-to-End Impact

TABLE IX
Production Impact Over 90-Day Evaluation

Metric	Baseline	InferTwin	Improvement
Unplanned downtime (hr/mo)	4.2	2.8	−34%
Mean throughput (tok/s)	3,820	4,240	+11%
Energy (kWh/M tokens)	1.71	1.57	−8.3%
Anomaly lead time (min)	2.1	28.4	+13.5×
Capacity planning time	2–4 weeks	<1 min	>1000×

X. Discussion

Principal findings: (1) Neural SDE world models achieve <4.2% throughput prediction error over 24 hours—sufficient for production planning. (2) Thermal cascades across racks cause 6–15% throughput loss invisible to per-node monitoring. (3) Causal counterfactual queries replace weeks of physical capacity planning with 50 ms simulations. (4) Offline RL in the world model produces policies that transfer to production with <5% sim-to-real gap. (5) Divergence-based anomaly detection provides 12–47 minute lead time before failures.

Limitations: (1) The Neural SDE requires >30 days of telemetry for initial training. (2) Rare catastrophic failures (UPS total failure, seismic events) lack sufficient training data. (3) Hardware architecture changes (new GPU generation) require re-training of affected twin levels.

Industry implications: InferTwin represents a paradigm shift from reactive monitoring to predictive infrastructure management. The world model serves as a shared mental model that aligns operations, capacity planning, and financial teams around a single source of truth.

Patent-protected innovations: (1) Neural SDE world model for joint inference + DC physical infrastructure simulation. (2) Causal counterfactual engine for capacity planning via do-calculus on learned dynamics. (3) Predictive anomaly detection via world model state divergence.

XI. Conclusion

We presented InferTwin, a continuous-time neural world model that unifies inference control plane optimization with datacenter infrastructure digital twinning. Our six contributions—Neural SDE formulation, hierarchical 4-level twin, causal counterfactual engine, DreamerV3-Infra offline RL, divergence-based anomaly detection, and DC foundation model transfer—collectively reduce unplanned downtime by 34%, improve throughput by 11%, and cut energy consumption by 8.3%. InferTwin transforms datacenter operations from reactive threshold monitoring to predictive, causal, simulation-driven management.

References

[1] NVIDIA Corp., "Nvidia data center gpu manager," 2025.
[2] R. T. Chen, Y. Rubanova, J. Bettencourt, and D. Duvenaud, "Neural ordinary differential equations," NeurIPS, 2018.
[3] X. Li, T.-K. L. Wong, R. T. Chen, and D. Duvenaud, "Scalable gradients for stochastic differential equations," AISTATS, 2020.
[4] D. Hafner, J. Pasukonis, J. Ba, and T. Lillicrap, "Mastering diverse domains through world models," arXiv:2301.04104, 2023.
[5] J. Schrittwieser et al., "Mastering atari, go, chess and shogi by planning with a learned model," Nature, 2020.
[6] M. Grieves and J. Vickers, "Digital twin: Manufacturing excellence through virtual factory replication," 2014.
[7] F. Tao et al., "Digital twin in industry: State-of-the-art," IEEE TII, 2018.
[8] M. Kiran et al., "Digital twins for data center infrastructure management," IEEE Cloud, 2022.
[9] H. Mao et al., "Learning scheduling algorithms for data processing clusters," SIGCOMM, 2019.
[10] A. Mirhoseini et al., "A graph placement methodology for fast chip design," Nature, 2021.
[11] L. Chen et al., "Decision transformer: Rl via sequence modeling," NeurIPS, 2021.
[12] J. Pearl, "Causality," 2009.
[13] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, "Proximal policy optimization algorithms," arXiv:1707.06347, 2017.
[14] J. Achiam et al., "Constrained policy optimization," ICML, 2017.
[15] J. Tobin et al., "Domain randomization for sim-to-real transfer," IROS, 2017.