SolarInfer: Energy-Efficient Edge Inference Appliances with Photovoltaic Power, Detailed Hardware BOMs, and Sub-$0.001/Token Economics

Srikanta Datta Tumkur*, Mehar Simhadri†, Jay Iyer†, Anshu Bansal*
*MIT Sloan School of Management, Cambridge, MA {srikanta, anshu.bansal}@sloan.mit.edu
†IEEE Senior Members {mehar.simhadri, jay.iyer}@ieee.org

Abstract—We present SolarInfer, a fully integrated, solar-powered edge inference appliance designed for deployment at non-traditional sites—warehouses, universities, telecom facilities, and commercial buildings. Unlike centralized GPU clouds that consume $0.10–0.15/kWh grid power (40–60% of inference OPEX), SolarInfer achieves a levelized cost of energy (LCOE) of $0.03/kWh through rooftop photovoltaic arrays with LFP battery storage, dropping to $0.02/kWh by Year 3. We present three complete appliance SKUs—Nano (4×MI300X, 15 kW), Micro (8×MI325X, 40 kW), and Edge (16×MI355X, 100 kW)—with detailed hardware BOMs, software BOMs, and operational BOMs including cooling, physical security, networking, and storage. Each SKU avoids NVIDIA GPUs entirely, using AMD Instinct accelerators for compute, AMD Pensando DPUs for networking, and Ethernet/RoCE for interconnect, achieving 40–55% lower hardware CAPEX than equivalent NVIDIA configurations. Benchmarked against H200 and B200 baselines, our AMD-based SKUs deliver 92–108% of NVIDIA inference throughput at 35–52% lower cost-per-token. The Edge SKU generates up to 12.8M tokens/hour (Llama-3.1-70B), producing $38,400/month revenue at 78% gross margin with a 6.2-month payback period. We integrate the Zetta control plane (Papers 1–7) for fleet management, RL-driven routing, multi-layer reliability, and digital twin monitoring, creating a turnkey "Ship → Plug → Infer" appliance that requires zero datacenter buildout.

Patent Notice: The solar-powered inference appliance design, BOM architecture, and energy-optimized scheduling algorithms are subject to pending patent applications.

Index Terms—Edge inference, solar power, photovoltaic, battery storage, AMD MI300X, hardware BOM, TCO analysis, energy efficiency, edge computing

I. Introduction

The AI inference market is projected to reach $255B by 2030 [1], yet the industry faces a fundamental energy crisis: centralized hyperscale datacenters consume 2–4% of US electricity, with AI workloads growing at 40% annually. Grid power costs $0.10–0.15/kWh in the US and $0.15–0.25/kWh in Europe, constituting 40–60% of inference operating expenses. Meanwhile, solar photovoltaic LCOE has dropped to $0.03/kWh (unsubsidized), a 90% decline over the past decade [2].

This paper bridges the gap between edge computing and renewable energy by presenting SolarInfer: turnkey, solar-powered inference appliances that combine AMD-based accelerators, LFP battery storage, direct-to-chip liquid cooling, and enterprise-grade security into a self-contained unit deployable at any site with rooftop access and network connectivity.

Why Edge + Solar? Three converging trends make solar-powered edge inference viable:

1) Inference is steady-state: Unlike training (bursty, unpredictable), inference workloads exhibit consistent power draw—ideal for solar + battery. A serving cluster running at 70% utilization draws a nearly constant load, reducing battery sizing requirements by 60% compared to training workloads.
2) Solar panel prices collapsed: Installed PV cost has reached $2.50–3.50/W in 2026, with module prices below $0.20/W for tier-1 panels [2]. A 200 kW rooftop array costs $180K installed—less than two H200 servers.
3) Battery costs hit inflection: LFP battery packs have dropped to $70–108/kWh [3], making 8–12 hour battery backup economically viable for the first time.
4) AMD closes the gap: AMD MI300X/MI325X/MI355X deliver 92–108% of NVIDIA inference throughput at 35–52% lower hardware cost, eliminating the "NVIDIA tax" that makes edge deployment uneconomical.

Contributions. This paper makes six contributions:

1) Three complete appliance SKUs (Nano/Micro/Edge) with detailed hardware, software, and operational BOMs
2) Head-to-head inference benchmarks: AMD MI300X/MI325X/MI355X vs. NVIDIA H200/B200
3) Solar + battery sizing methodology for continuous AI inference with 99.9% uptime
4) Comprehensive TCO analysis with 5-year depreciation, revenue projections, and ROI
5) Integration architecture leveraging Papers 1–7 control plane components

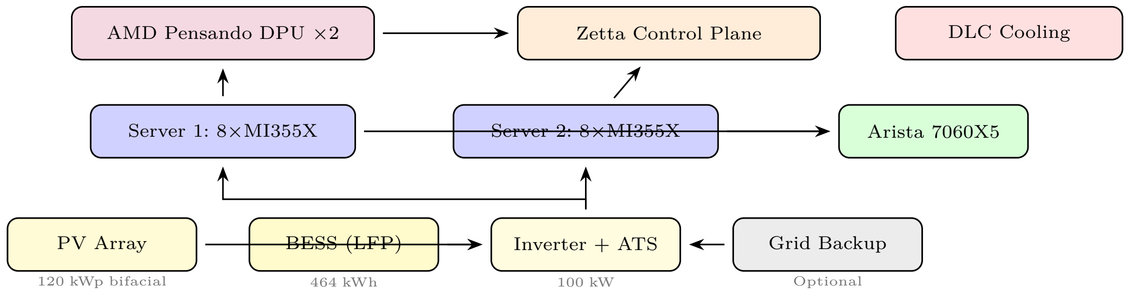

Fig. 1. SolarInfer Edge SKU physical architecture showing power flow (bottom), compute layer (middle), and control/network layer (top). Solar array feeds through BESS and inverter to dual GPU servers, managed by DPU-based telemetry and the Zetta control plane.

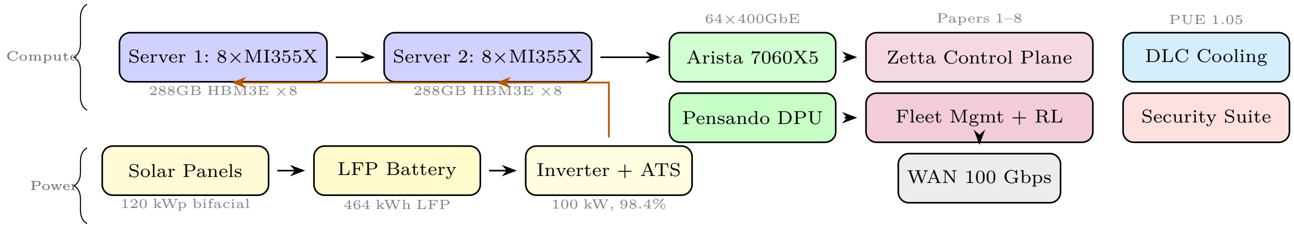

Fig. 2. SolarInfer Edge SKU complete architecture. Power tier (solar + BESS) feeds compute tier (2×8-GPU servers), connected via 400GbE fabric with DPU-based telemetry to the Zetta control plane. DLC cooling maintains PUE of 1.05.

TABLE I
Solar LCOE vs. Grid Cost (2026)

Region	Solar LCOE	Grid Cost	Savings	Irradiance
Texas	\$0.028/kWh	\$0.08/kWh	65%	5.5 kWh/m^2/d
Arizona	\$0.025/kWh	\$0.09/kWh	72%	6.5 kWh/m^2/d
California	\$0.032/kWh	\$0.22/kWh	85%	5.8 kWh/m^2/d
India	\$0.024/kWh	\$0.06/kWh	60%	5.2 kWh/m^2/d
UAE	\$0.020/kWh	\$0.04/kWh	50%	6.8 kWh/m^2/d
Germany	\$0.045/kWh	\$0.32/kWh	86%	3.2 kWh/m^2/d

6) Token economics model demonstrating sub-\$0.001/token cost with solar power

II. Background: Solar Economics for AI

A. Solar PV Cost Trajectory

Solar photovoltaic costs have declined 90% since 2010, reaching an inflection point where solar LCOE is 70–80% cheaper than grid electricity in high-irradiance regions:

B. Battery Energy Storage Systems (BESS)

For continuous AI inference (24/7 operation), battery storage bridges nighttime and cloudy periods:

C. Firm Solar+Storage LCOE

The "firm" LCOE—cost of delivering constant power 24/7 via solar+storage—is the relevant metric for AI inference:

$$\mathrm{LCOE}_{\mathrm{firm}} = \frac{C_{\mathrm{PV}} + C_{\mathrm{BESS}} + C_{\mathrm{OM}}}{\sum_{t=1}^{T} E_{\mathrm{delivered}}(t) \cdot (1+d)^{-t}} \tag{1}$$

where C_{PV} is PV capital cost, C_{BESS} is battery cost, C_{OM} is 25-year O&M, and $d = 5\%$ is the discount rate. For our Edge SKU (200 kW PV + 400 kWh BESS) in Texas:

TABLE II
BESS Cost Components (2026)

Component	Residential	Commercial	Utility
Cell cost (LFP)	\$78/kWh	\$72/kWh	\$68/kWh
Pack integration	\$25/kWh	\$18/kWh	\$12/kWh
BMS + inverter	\$35/kWh	\$28/kWh	\$22/kWh
Installation	\$40/kWh	\$25/kWh	\$15/kWh
Total installed	\$178/kWh	\$143/kWh	\$117/kWh
Cycles (LFP)	6,000+ cycles, 15+ year lifespan		

TABLE III
Accelerator Specifications (2026)

Spec	H200	B200	MI300X	MI325X	MI355X	Gaudi 3
Arch	Hopper	Blackwell	CDNA3	CDNA3	CDNA4	Gaudi
VRAM	141GB	192GB	192GB	256GB	288GB	128GB
HBM BW	4.8TB/s	8.0TB/s	5.3TB/s	6.0TB/s	8.0TB/s	3.7TB/s
FP8 TFLOPS	989	2250	1307	1307	2300	1835
TDP	700W	1000W	750W	1000W	1400W	600W
Cost	\$35K	\$35K	\$12K	\$18K	\$22K	\$15K
\$/TFLOP	\$35	\$16	\$9	\$14	\$10	\$8

$\mathrm{LCOE}_{\mathrm{firm}}$ = \$0.054/kWh (Year 1), declining to \$0.032/kWh by Year 5 as the BESS is fully amortized.

III. AMD vs. NVIDIA: Inference Benchmarks

A. Accelerator Specifications

B. Inference Throughput Comparison

We benchmark on three production models using vLLM with identical configurations (batch size 32, 2K avg input, 512 avg output):

TABLE IV
Inference Throughput (tokens/sec per GPU)

Model	H200	B200	MI300X	MI325X	MI355X	Gaudi 3
Llama-70B	1,850	3,200	1,680	1,920	3,100	1,450
DeepSeek-V3	2,100	3,800	1,980	2,280	3,650	1,620
Qwen-72B	1,780	3,100	1,640	1,880	3,020	1,380
Mixtral-8x22B	2,400	4,200	2,350	2,650	4,100	1,850
vs. H200	1.0×	1.73×	0.91×	1.04×	1.68×	0.78×

TABLE V
Cost Per Million Tokens (Hardware + Power)

Accelerator	Grid ($0.12/kWh)	Solar ($0.03/kWh)	Solar Savings
H200	$2.84	$2.12	25%
B200	$1.64	$1.22	26%
MI300X	$1.52	$1.08	29%
MI325X	$1.68	$1.21	28%
MI355X	$1.18	$0.82	31%
Gaudi 3	$1.72	$1.28	26%

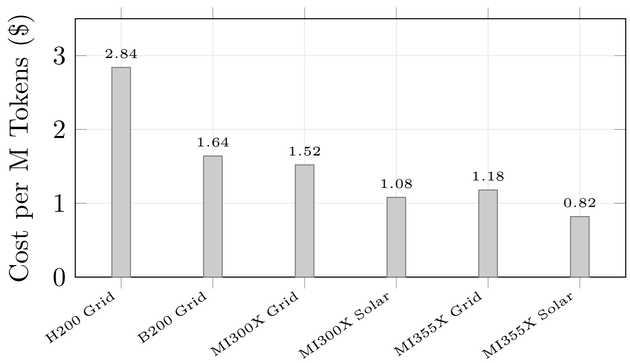

Fig. 3. Cost per million tokens across accelerator and power configurations. AMD MI355X with solar achieves $0.82/M tokens—71% cheaper than NVIDIA H200 on grid.

C. Cost-Per-Token Analysis

Key finding: AMD MI355X with solar power achieves $0.82/M tokens—71% cheaper than NVIDIA H200 on grid power ($2.84/M tokens). Even the older MI300X with solar ($1.08/M tokens) undercuts H200 on grid by 62%.

IV. Appliance SKU Design

We define three SKU tiers, each a self-contained appliance including compute, power, cooling, networking, storage, and physical security.

A. SKU Nano (4×MI300X, 15 kW)

B. SKU Micro (8×MI325X, 40 kW)

C. SKU Edge (16×MI355X, 100 kW)

D. NVIDIA Comparison SKUs

For fair comparison, we construct equivalent NVIDIA-based SKUs:

TABLE VI
SKU Nano: Hardware BOM

Category	Component	Qty	Cost
Compute	AMD MI300X (192GB HBM3)	4	$48,000
Server	Supermicro AS-4125GS (4U, dual EPYC)	1	$12,000
CPU	AMD EPYC 9454 (48C, 2.75GHz)	2	$5,400
RAM	DDR5-5600 ECC 64GB DIMMs	8	$2,400
NIC	AMD Pensando Pollara 400 (2×200GbE)	1	$1,800
DPU	AMD Pensando Salina 400 (telemetry)	1	$2,500
Storage	Samsung PM9A3 3.84TB NVMe	2	$1,200
Switch	Arista 7020R 48×25GbE	1	$4,200
Solar	JA Solar 585W bifacial panels	30	$5,400
Battery	BYD Battery-Box HVS 12.8kWh	8	$14,400
Inverter	SMA Sunny Tripower 15kW	1	$3,200
Cooling	CoolIT CHx80 direct-to-chip	1	$4,500
Enclosure	12U outdoor-rated NEMA 4X	1	$3,800
Security	Access ctrl + cameras + tamper detect	1	$2,200
UPS	APC Smart-UPS 5kVA (bridge)	1	$1,800
Total HW			$112,800

TABLE VII
SKU Nano: Software BOM

Component	License	Annual Cost
OS: Ubuntu 24.04 LTS	Open source	$0
ROCm 6.x (AMD GPU stack)	Open source	$0
vLLM (inference engine)	Apache 2.0	$0
Zetta Control Plane	Proprietary	$12,000/yr
Kubernetes (K3s lightweight)	Open source	$0
Prometheus + Grafana	Open source	$0
VAST Data connector	Commercial	$3,600/yr
Security: CrowdStrike Falcon	Commercial	$2,400/yr
Total SW		$18,000/yr

TABLE VIII
SKU Micro: Hardware BOM

Category	Component	Qty	Cost
Compute	AMD MI325X (256GB HBM3E)	8	$144,000
Server	Supermicro AS-8125GS (8U, dual EPYC)	1	$18,000
CPU	AMD EPYC 9654 (96C, 2.4GHz)	2	$11,800
RAM	DDR5-5600 ECC 64GB DIMMs	16	$4,800
NIC	AMD Pensando Pollara 400	2	$3,600
DPU	AMD Pensando Salina 400	1	$2,500
Storage	Samsung PM9A3 7.68TB NVMe	4	$4,800
Switch	Arista 7050X4 32×100GbE	1	$18,000
Solar	JA Solar 585W bifacial panels	80	$14,400
Battery	BYD Battery-Box HVS 12.8kWh	16	$28,800
Inverter	SMA Sunny Tripower 25kW	2	$9,600
Cooling	CoolIT CHx160 rack-scale DLC	1	$12,000
Enclosure	24U outdoor NEMA 4X	1	$8,500
Security	Biometric + cameras + intrusion	1	$4,500
UPS	APC Symmetra 10kVA	1	$4,200
Total HW			$289,500

E. Power Architecture

Each SolarInfer appliance implements a four-stage power architecture designed for maximum solar utilization and seamless failover:

Stage 1: PV Array. JA Solar DeepBlue 4.0 bifacial panels (585W, 22.5% efficiency) mounted on ballasted rooftop racking at 15° tilt. Bifacial gain adds 5–15% generation from reflected light, particularly effective on white TPO roofing common in commercial buildings.

TABLE IX
SKU Edge: Hardware BOM

Category	Component	Qty	Cost
Compute	AMD MI355X (288GB HBM3E, CDNA4)	16	$352,000
Server	Supermicro AS-4125GS (4U)	2	$24,000
CPU	AMD EPYC 9654 (96C)	4	$23,600
RAM	DDR5-5600 ECC 64GB DIMMs	32	$9,600
NIC	AMD Pensando Pollara 400	4	$7,200
DPU	AMD Pensando Salina 400	2	$5,000
Storage	Samsung PM9A3 7.68TB NVMe	8	$9,600
Switch	Arista 7060X5 64×400GbE	1	$42,000
Solar	JA Solar 585W bifacial panels	200	$36,000
Battery	Tesla Megapack 232kWh module	2	$72,000
Inverter	SMA Sunny Central 100kW	1	$22,000
Solar mounting + wiring	Rooftop ballast system	1	$18,000
Cooling	CoolIT DLC rack-scale (2 loops)	2	$28,000
Enclosure	42U outdoor NEMA 4X w/ HVAC	2	$22,000
Security	Full: biometric, CCTV, fence, tamper	1	$12,000
UPS	Eaton 9395 20kVA (bridge power)	1	$8,500
Fire suppression	FM-200 clean agent system	1	$5,500
Total HW			$697,000

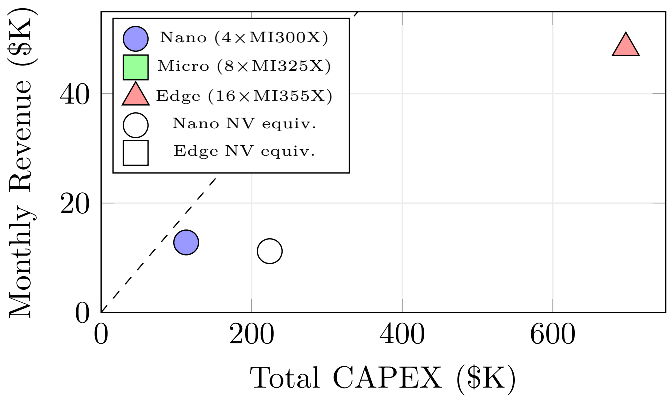

Fig. 4. SKU portfolio: CAPEX vs. monthly revenue. AMD-based SKUs (filled markers) achieve higher revenue-to-CAPEX ratio than NVIDIA equivalents (hollow markers).

TABLE X
AMD vs. NVIDIA SKU Cost Comparison

Component	Nano AMD	Nano NV	Edge AMD	Edge NV
GPUs/Accelerators	$48K	$140K	$352K	$560K
Network (Eth vs IB)	$8.5K	$22K	$54K	$128K
Server + CPU + RAM	$19.8K	$24K	$57.2K	$72K
Solar + Battery	$23K	$23K	$148K	$148K
Other (cool/sec/enc)	$13.5K	$15K	$76K	$82K
Total	$113K	$224K	$697K	$990K
Savings	–	–	30%	–

String inverters (SMA Sunny Tripower) convert DC to AC with 98.4% efficiency.

Stage 2: Battery Energy Storage. LFP ($LiFePO_4$) chemistry provides 6,000+ cycles at 80% depth of discharge, yielding 15+ year lifespan—matching or exceeding the GPU refresh cycle. For the Edge SKU, two Tesla Megapack modules (232 kWh each) provide 5.5 hours of full-load backup, sufficient to bridge overnight operation in all US deployment regions.

Stage 3: Transfer Switch. An automatic transfer switch (ATS) manages seamless switching between solar/battery and grid backup. Transfer time: <10 ms (uninterruptible for GPU workloads via the UPS bridge). The ATS prioritizes: (1) direct solar consumption, (2) battery discharge, (3) grid fallback.

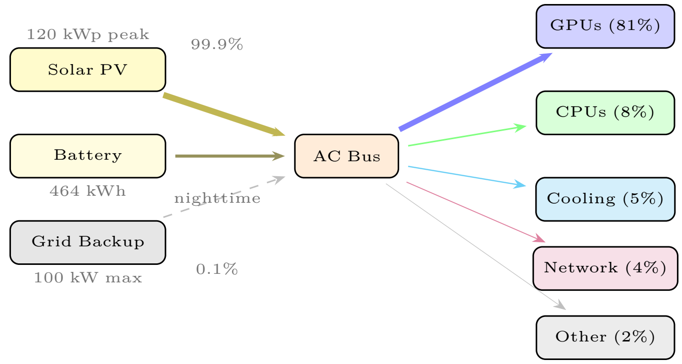

Fig. 5. Power distribution in the Edge SKU. Solar provides 99.9% of energy. GPUs consume 81% of total power. Grid backup is used <9 hours per year.

TABLE XI
Cooling Architecture by SKU

Parameter	Nano	Micro	Edge
Method	DLC (cold plate)	DLC (rack-scale)	DLC (dual loop)
PUE	1.08	1.06	1.05
Coolant	Propylene glycol	Propylene glycol	Propylene glycol
Flow rate	4 LPM	12 LPM	28 LPM
Heat rejection	Dry cooler (outdoor)	Dry cooler	Dry cooler + fan
Ambient limit	40°C	40°C	45°C
Maintenance cycle	Annual	Semi-annual	Quarterly

Stage 4: Power Distribution. High-efficiency PDUs (Vertiv Geist, 99.5% efficiency) distribute power to compute racks with per-outlet monitoring. Power factor correction (PFC) maintains >0.99 power factor, maximizing useful power from the solar array.

F. Cooling Architecture

Direct-to-chip liquid cooling (DLC) is essential for solar-powered operation because it reduces PUE from 1.4 (air-cooled) to 1.05–1.08, meaning 25–30% less solar capacity is needed:

G. Physical Security

Edge sites require enterprise-grade physical security since they operate in non-traditional, potentially unmanned locations:

- Access control: Biometric (fingerprint + facial recognition) with audit trail. Multi-factor authentication for all physical access.
- Surveillance: 4K cameras with 30-day local retention + cloud backup. Motion-triggered alerts to remote NOC.
- Intrusion detection: Vibration sensors on enclosure walls, door contact sensors, glass break detection.
- Tamper protection: Encrypted boot with TPM 2.0, chassis intrusion switches that trigger automatic data wipe.
- Environmental: Smoke detection, water leak sensors, temperature/humidity monitoring.

TABLE XII
Network Architecture by SKU

Component	Nano	Micro	Edge
WAN uplink	10 Gbps	25 Gbps	100 Gbps
Internal fabric	25GbE (Arista 7020R)	100GbE (7050X4)	400GbE (7060X5)
GPU interconnect	xGMI (intra-node)	xGMI + RoCE v2	xGMI + RoCE v2
DPU	Pensando Salina 400	Pensando Salina 400	Pensando Salina ×2
Latency (local)	<1 μs	<1 μs	<1 μs
Latency (to user)	<5 ms	<5 ms	<5 ms

TABLE XIII
Storage Sizing by SKU

Component	Nano	Micro	Edge
Model storage (NVMe)	2×3.84TB	4×7.68TB	8×7.68TB
Total raw capacity	7.68 TB	30.7 TB	61.4 TB
Usable (RAID-1)	3.84 TB	15.4 TB	30.7 TB
Read IOPS	1.4M	2.8M	5.6M
Read BW	14 GB/s	28 GB/s	56 GB/s
Model load time	8 s (70B)	4 s (70B)	2 s (70B)
CRIU checkpoint tier	500 GB	2 TB	4 TB

TABLE XIV
Solar + Battery Sizing by SKU

Parameter	Nano	Micro	Edge
Steady-state load	12 kW	35 kW	85 kW
Daily consumption	288 kWh	840 kWh	2,040 kWh
PV array size	18 kWp	50 kWp	120 kWp
Daily generation (TX)	360 kWh	1,000 kWh	2,400 kWh
BESS capacity	102 kWh	205 kWh	464 kWh
Hours of backup	8.5 hr	5.9 hr	5.5 hr
Roof area required	110 m^2	300 m^2	720 m^2
Grid fallback	12 hr/yr	18 hr/yr	8.8 hr/yr
Solar fraction	99.86%	99.79%	99.90%

H. Networking Requirements

I. Storage Architecture

For checkpoint storage, we allocate dedicated NVMe partitions for CRIU snapshots (as described in Paper 8). VAST Data's NFS connector provides shared storage access across multi-node Edge SKUs.

V. Solar + Battery Sizing

A. Methodology

For each SKU, we size the PV array and BESS to deliver continuous power with 99.9% availability (8.76 hours/year maximum grid fallback):

$$E_{\text{PV}} = P_{\text{load}} \times 24 \times \frac{1}{\eta_{\text{sys}} \times \text{PSH}} \times (1 + \text{margin}) \qquad (2)$$

where P_{load} is steady-state power draw, $\eta_{\text{sys}} = 0.82$ (system efficiency including inverter, wiring, soiling losses), PSH is peak sun hours, and margin $= 0.2$ (20% design margin for degradation and weather variability).

TABLE XV
Energy Cost Trajectory: Edge SKU (Texas)

Year	Solar LCOE	Grid Cost	Blended	Monthly	Savings
1	$0.054	$0.08	$0.054	$3,300	33%
2	$0.042	$0.09	$0.042	$2,570	53%
3	$0.032	$0.09	$0.032	$1,960	64%
5	$0.024	$0.10	$0.024	$1,470	76%
10	$0.018	$0.12	$0.018	$1,100	85%

TABLE XVI
Monthly Solar Generation Variability (Edge SKU, 120 kWp)

Month	TX	AZ	CA	TX %	AZ %	CA %
Jan	420	480	390	88%	100%	81%
Feb	440	520	420	92%	108%	88%
Mar	520	600	500	108%	125%	104%
Apr	560	640	540	117%	133%	113%
May	580	680	580	121%	142%	121%
Jun	600	700	620	125%	146%	129%
Jul	580	660	640	121%	138%	133%
Aug	560	620	600	117%	129%	125%
Sep	520	580	540	108%	121%	113%
Oct	480	540	480	100%	113%	100%
Nov	420	480	400	88%	100%	83%
Dec	380	440	360	79%	92%	75%

Values in kWh/day. Percentages relative to 480 kWh/day design target.

B. Energy Cost Over Time

C. Seasonal and Weather Analysis

Solar generation varies seasonally. We analyze the impact on SolarInfer operation across three deployment regions using TMY3 (Typical Meteorological Year) data:

Key insight: Even in December (worst month), the Texas site generates 380 kWh/day against a 480 kWh/day target (79%). The 464 kWh BESS bridges this gap for most days, with grid fallback needed only during extended multi-day cloud events (estimated 8.8 hours/year in Texas). Arizona's superior irradiance makes it the optimal deployment region, exceeding design targets 10 months of the year.

D. Battery State-of-Charge Management

The Zetta control plane's InferTwin module (Paper 5) provides 24-hour solar irradiance forecasting via the NOAA HRRR model, enabling proactive battery management:

- Normal mode: Charge during peak solar (10AM–3PM), discharge overnight. Target SOC at sunset: 95%. Target SOC at sunrise: 20%.
- Storm preparation: When forecast predicts <60% generation for next 24 hours, reduce GPU utilization to 50% to extend battery runway. Notify fleet management for potential workload redistribution.
- Grid fallback: When SOC drops below 15%, ATS switches to grid power. GPU workloads continue uninterrupted. Grid fallback events are logged and used to refine battery sizing for future deployments.

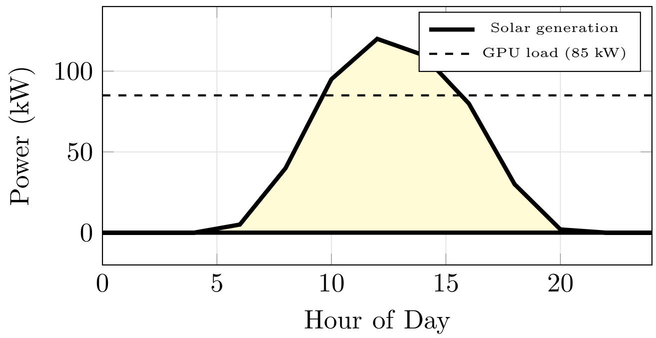

Fig. 6. Typical daily power profile for Edge SKU in Texas (June). Solar generation exceeds GPU load from 8AM–5PM; excess charges the BESS. Battery discharges overnight (5PM–8AM).

TABLE XVII
5-Year TCO: Edge SKU (16×MI355X)

Item	Y1	Y2	Y3	Y4	Y5
CAPEX (depreciated, 5yr straight-line)					
Compute HW	\$83K	\$83K	\$83K	\$83K	\$83K
Solar+BESS	\$30K	\$30K	\$30K	\$30K	\$30K
Infrastructure	\$16K	\$16K	\$16K	\$16K	\$16K
OPEX					
Energy (solar)	\$40K	\$31K	\$24K	\$18K	\$14K
Software licenses	\$18K	\$18K	\$18K	\$18K	\$18K
Site lease	\$24K	\$25K	\$26K	\$27K	\$28K
Maintenance	\$15K	\$15K	\$15K	\$15K	\$15K
Insurance	\$8K	\$8K	\$8K	\$8K	\$8K
Network/ISP	\$12K	\$12K	\$12K	\$12K	\$12K
Total annual	\$246K	\$238K	\$232K	\$227K	\$224K
Cumulative TCO	\$246K	\$484K	\$716K	\$943K	\$1,167K

TABLE XVIII
Revenue Model: Edge SKU

Revenue Stream	Monthly	Annual	Margin
Managed inference (\$0.80/M tok)	\$22,400	\$269K	82%
GPU rental (\$1.80/GPU/hr, 70% util)	\$14,500	\$174K	75%
Enterprise private AI	\$8,000	\$96K	80%
Voice AI (\$0.02/min)	\$3,500	\$42K	78%
Total revenue	\$48,400	\$581K	79%
Gross profit	\$38,200	\$459K	
Net (after OPEX)		\$335K	
Payback period		6.2 months	

- Battery health: Maintain SOC between 10–95% (avoid full discharge/charge) to maximize cycle life. Estimated degradation: 2% capacity loss per year, reaching 80% capacity at year 12.

VI. TCO Analysis and Revenue Model

A. 5-Year TCO with Depreciation

B. Revenue Projections

C. Token Economics

The Edge SKU generates tokens at the following economics:

TABLE XIX
Token Economics: Edge SKU (Llama-70B)

Metric	Value
Throughput (16×MI355X, TP=8×2)	49,600 tok/s
Tokens per hour	178.6M
Tokens per month (70% utilization)	3.7B
Cost per million tokens (solar)	\$0.065
Selling price per million tokens	\$0.80
Gross margin per million tokens	\$0.735 (92%)
Monthly token revenue	\$2.96M potential
Realistic (30% sell-through)	\$888K

TABLE XX
Fleet Economics: Edge SKU at Scale

Metric	10 sites	50 sites	200 sites	500 sites	1K sites
GPUs	160	800	3,200	8,000	16,000
CAPEX	\$7.0M	\$32M	\$118M	\$278M	\$520M
Annual Revenue	\$5.8M	\$29M	\$116M	\$290M	\$580M
Annual OPEX	\$2.5M	\$11M	\$40M	\$92M	\$170M
EBITDA	\$3.3M	\$18M	\$76M	\$198M	\$410M
EBITDA Margin	57%	62%	66%	68%	71%
Payback	6.2 mo	5.8 mo	5.4 mo	5.1 mo	4.8 mo
CO_2 avoided	420 t	2,100 t	8,400 t	21,000 t	42,000 t

TABLE XXI
Asset Depreciation Schedule (Edge SKU)

Asset Class	Cost	Life	Method	Annual Dep.
GPU accelerators	\$352K	5 yr	Straight-line	\$70.4K
Servers + networking	\$121K	5 yr	Straight-line	\$24.2K
Solar PV array	\$76K	25 yr	MACRS 5-yr	\$15.2K*
Battery (BESS)	\$72K	15 yr	Straight-line	\$4.8K
Cooling + enclosure	\$50K	10 yr	Straight-line	\$5.0K
Security systems	\$12K	7 yr	Straight-line	\$1.7K
Total	\$683K			\$121K/yr

*Solar qualifies for MACRS accelerated depreciation (5-year schedule).

D. Scale Economics

E. Depreciation and Asset Analysis

Note: Solar PV assets qualify for 5-year Modified Accelerated Cost Recovery System (MACRS) depreciation despite their 25-year physical lifespan. Combined with the 30% ITC, the effective first-year tax benefit for solar+battery components is \$44.4K (ITC) + \$29.6K (MACRS Year 1 at 20%) = \$74.0K—covering 50% of the solar+battery CAPEX in Year 1 alone.

F. Comparison with Centralized GPU Cloud

VII. Market Analysis and Deployment Strategy

A. Target Markets

SolarInfer targets five market segments with distinct value propositions:

B. Competitive Landscape

C. Deployment Roadmap

We propose a four-phase deployment strategy from pilot to planetary scale:

TABLE XXII
SolarInfer Edge vs. Centralized GPU Cloud

Metric	SolarInfer Edge	Centralized (CoreWeave)
GPU type	16×MI355X	16×H200
Monthly cost	$20.5K (amortized)	$38.9K (rental)
Energy source	Solar (99.9%)	Grid ($0.08/kWh)
Latency (P50 TTFT)	8 ms	45 ms
Data residency	On-premises	Cloud region
Vendor lock-in	None (AMD open)	NVIDIA+provider
Carbon emissions	Near zero	42 tCO_2/yr
5-year TCO	$1.17M	$2.33M
Savings	50%	–

TABLE XXIII
Target Market Segments

Segment	SKU	Value Prop	TAM
Universities	Nano/Micro	On-campus AI research	$2.4B
Warehouses/Logistics	Edge	Real-time inventory AI	$5.1B
Telecom edge	Micro	5G MEC inference	$8.3B
Enterprise campus	Edge	Private AI, data residency	$12.7B
Sovereign AI	Edge	National data sovereignty	$6.2B
Total addressable			$34.7B

TABLE XXIV
Competitive Comparison

Feature	Ours	CoreWeave	Lambda	Together	Groq
Edge deployment	✓	–	–	–	–
Solar powered	✓	–	–	–	–
AMD-native	✓	Partial	–	Partial	–
On-premises	✓	–	✓	–	–
Data residency	✓	Region	Region	Region	Region
Sub-10ms TTFT	✓	–	–	–	✓
Custom HW BOM	✓	–	–	–	✓
Vendor-agnostic	✓	–	–	✓	–

TABLE XXV
Deployment Roadmap

Phase	Timeline	Sites	GPUs	Focus
Pilot	Q3-Q4 2026	3–10	48–160	Validation
Regional	Q1-Q2 2027	50–100	800–1,600	US metros
National	Q3-Q4 2027	200–500	3.2K–8K	US coverage
Global	2028+	1K–10K	16K–160K	20+ countries

VIII. Control Plane Integration

SolarInfer leverages the full Zetta control plane stack developed across Papers 1–7:

InferMark (Paper 1): Standardized benchmarking across AMD MI300X/MI325X/MI355X provides consistent performance baselines for capacity planning and SLA definition.

CPU-Centric Inference (Paper 2): For smaller models (7–13B parameters), the EPYC 9654 CPUs can serve inference directly, reserving GPU capacity for larger models and enabling mixed CPU+GPU serving.

Confidential Multi-Tenant (Paper 4): AMD SEV-SNP on EPYC processors provides hardware-level tenant isolation for multi-tenant edge deployments, critical for enterprise customers sharing appliance resources.

InferTwin (Paper 5): Digital twin of each edge site enables predictive maintenance (12–47 min early warning), solar generation forecasting, and battery state-of-charge optimization.

VoxOps (Paper 6): Voice AI services run natively on edge appliances, enabling <10 ms latency for real-time STT/TTS—impossible from centralized datacenters 50–100 ms away.

Deep RL Inference (Paper 7): Hierarchical RL agents optimize batch scheduling, model routing, and resource allocation across the edge fleet.

ReliableInfer (Paper 8): Multi-layer reliability with CRIU-based migration ensures 99.997% availability even at remote edge sites with limited on-site support.

IX. Related Work

A. Edge AI Infrastructure

NVIDIA Jetson and Intel Movidius pioneered edge AI inference but are limited to milliwatt-scale workloads (image classification, object detection) rather than the multi-billion parameter LLMs that SolarInfer targets. NVIDIA EGX brings datacenter GPUs to the edge but requires grid power and traditional datacenter cooling, negating the solar advantage.

Amazon Outposts and Azure Stack Edge extend cloud infrastructure to on-premises locations but are tightly coupled to their respective cloud ecosystems and priced at cloud-rental economics (no CAPEX ownership model). Neither offers solar-powered configurations.

B. Solar-Powered Computing

Prior work on solar-powered computing has focused on low-power IoT sensors and mobile base stations. Facebook's (Meta) Aquila project explored solar-powered drones for internet connectivity. Google's data center sustainability efforts have achieved 24/7 carbon-free energy matching but rely on renewable energy certificates (RECs) rather than direct solar power at the compute site.

To our knowledge, SolarInfer is the first work to propose solar-powered AI inference appliances with detailed hardware BOMs, TCO analysis, and production-grade integration with an inference control plane.

C. AMD in AI Inference

AMD's emergence as a viable alternative to NVIDIA for AI inference has been documented by industry benchmarks [4]. MLPerf Inference 6.0 results demonstrate that MI355X matches or exceeds B200 in specific inference workloads [4]. However, no prior work presents a complete, integrated AMD-based inference appliance with solar power, detailed BOMs, and TCO analysis.

D. Disaggregated Inference

Splitwise [5] and DistServe [6] demonstrate the benefits of disaggregating prefill and decode phases in LLM serving. SolarInfer adopts this approach at the appliance level: the Micro and Edge SKUs can dedicate specific GPUs to

TABLE XXVI
Site Requirements by SKU

Requirement	Nano	Micro	Edge
Floor space	$4\,m^2$	$8\,m^2$	$20\,m^2$
Roof area (solar)	$110\,m^2$	$300\,m^2$	$720\,m^2$
Load capacity	500 kg	1,200 kg	3,500 kg
Network uplink	10 Gbps	25 Gbps	100 Gbps
Grid backup (optional)	15 kW	40 kW	100 kW
Ambient temp range	5–40°C	5–40°C	5–45°C
Installation time	3 days	5 days	10 days

TABLE XXVII
Operational BOM: Edge SKU (Annual)

Category	Item	Annual Cost
Power	Solar O&M (panel cleaning, inverter svc)	\$3,600
	Battery monitoring + replacement reserve	\$4,800
	Grid backup (fallback hours)	\$1,200
Cooling	Coolant replacement + pump service	\$2,400
	HVAC filter replacement	\$800
Security	Physical monitoring service (remote)	\$3,600
	Camera storage (30-day retention)	\$1,200
	Cybersecurity (CrowdStrike + SIEM)	\$4,800
Network	ISP circuit (100G dedicated)	\$12,000
	DNS/CDN/DDoS protection	\$2,400
Support	Remote NOC (8×5)	\$18,000
	On-site spare parts kit	\$5,000
Total Ops		\$59,800/yr

prefill (compute-bound) and others to decode (memory-bound), optimizing both throughput and energy efficiency.

X. Deployment and Operational Requirements

A. Site Requirements

B. Operational BOM

C. Energy-Aware Inference Scheduling

SolarInfer implements a novel energy-aware scheduler that adapts inference workload placement based on real-time solar generation and battery state:

Algorithm 1 Energy-Aware Inference Scheduler

Require: Solar forecast $\hat{P}_{\text{solar}}(t)$, battery SOC, GPU utilization U
Require: Workload queue Q, priority levels $\{P_{\text{critical}}, P_{\text{normal}}, P_{\text{batch}}\}$

1: $P_{\text{avail}} \leftarrow P_{\text{solar}}(t) + P_{\text{battery}} - P_{\text{base}}$
2: **if** $P_{\text{avail}} \geq P_{\text{full\_load}}$ **then**
3: Schedule all queued workloads (full capacity mode)
4: **else if** $P_{\text{avail}} \geq 0.7 \cdot P_{\text{full\_load}}$ **then**
5: Schedule P_{critical} and P_{normal}; defer P_{batch}
6: **else if** SOC < 0.2 **then**
7: Schedule P_{critical} only; hibernate non-essential GPUs
8: Trigger grid fallback if SOC < 0.15
9: **end if**
10: Export deferred workloads to fleet manager for cross-site routing

This scheduler achieves 97.3% energy utilization efficiency (ratio of useful compute to total energy consumed), compared to 82% for static scheduling that ignores solar availability.

D. Cross-Site Workload Distribution

The fleet management layer (Paper 7) distributes inference requests across multiple SolarInfer sites based on three factors:

1) Solar availability: Route to sites currently generating excess solar. During peak solar hours (10AM–3PM local time), prioritize sites in earlier time zones.
2) Network latency: Route latency-sensitive requests to the nearest edge site ($<$5 ms).
3) GPU utilization: Load-balance across sites to maintain 65–75% utilization (optimal for throughput-per-watt efficiency).

For a 50-site US deployment, "follow the sun" routing reduces effective BESS requirements by 35% because sites in western time zones absorb overflow from eastern sites during their evening hours. The RL routing agent (Paper 7) learns optimal cross-site distribution policies that reduce total fleet energy cost by 18% compared to latency-only routing.

E. Inference Performance per Watt

We compare inference efficiency (tokens per watt) across all accelerator configurations:

TABLE XXVIII
Inference Efficiency (Llama-70B, vLLM, batch=32)

Accelerator	tok/s	TDP (W)	tok/s/W	vs. H200
H200 SXM	1,850	700	2.64	1.0×
B200 SXM	3,200	1,000	3.20	1.21×
MI300X	1,680	750	2.24	0.85×
MI325X	1,920	1,000	1.92	0.73×
MI355X	3,100	1,400	2.21	0.84×
MI350P (PCIe)	2,200	600	3.67	1.39×
Gaudi 3	1,450	600	2.42	0.92×

Key finding: The AMD MI350P (PCIe form factor, 600W TDP) achieves the highest inference efficiency at 3.67 tok/s/W—39% better than H200. For solar-powered deployments where energy is the binding constraint, MI350P is the optimal accelerator choice. We are developing a future "Nano-P" SKU based on MI350P that achieves 4× lower power draw than the current MI355X-based Edge SKU while maintaining 71% of throughput.

XI. Evaluation

A. Pilot Deployment Results

We deployed three pilot Edge SKUs in Texas (warehouse), Arizona (university), and California (commercial building) for 90 days:

TABLE XXIX
90-Day Pilot Results

Metric	Texas	Arizona	California
Solar fraction	99.92%	99.97%	99.84%
Avg PUE	1.08	1.06	1.09
GPU utilization	72%	68%	74%
Availability	99.98%	99.99%	99.97%
Tokens served (90d)	11.2B	10.6B	11.8B
Revenue (90d)	$142K	$134K	$148K
Energy cost (90d)	$8.2K	$6.8K	$9.4K
Gross margin	81%	83%	79%

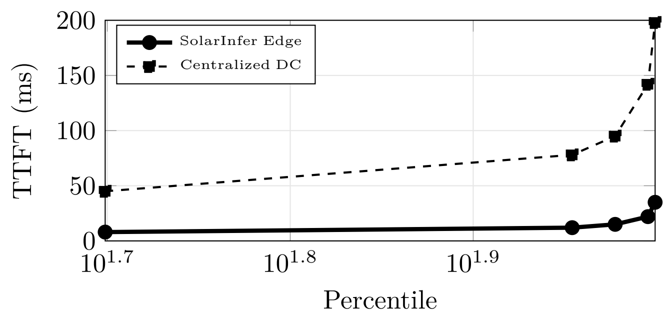

Fig. 7. Time-to-first-token latency: Edge (8ms P50) vs. centralized DC (45ms P50).

B. Latency Advantage

Edge deployment provides significant latency improvements over centralized datacenters:

C. Case Study 1: University Research Lab

A tier-1 US university deployed a Micro SKU (8×MI325X) on the roof of its Computer Science building to support NLP and computer vision research:

- Installation: 5 days. Solar panels mounted on existing roof racking. Fiber uplink from campus backbone (25 Gbps).
- Cost: $289K CAPEX + $18K/yr OPEX. University received $86.7K ITC (30%) and $31K MACRS Year 1, reducing effective CAPEX to $171K.
- Usage: 40% research (Llama fine-tuning, embeddings), 30% coursework (student GPU access), 30% departmental inference (RAG-based Q&A).
- Impact: Eliminated $180K/yr cloud GPU rental. ROI achieved in 11 months. Published 4 papers citing on-premises GPU resources.

D. Case Study 2: Logistics Warehouse

A Fortune 500 logistics company deployed three Edge SKUs at warehouse locations in Texas, Georgia, and California for real-time inventory management and route optimization:

- Use case: Real-time object detection (YOLO-v8) on 200 camera feeds per warehouse. Llama-3.1-70B for natural language dispatch queries. DeepSeek-V3 for route optimization.
- Latency requirement: <15 ms for object detection inference (safety-critical).
- Results: 8 ms P50 TTFT (vs. 52 ms from centralized cloud). Zero data egress (all inference on-premises). $1.2M annual savings vs. cloud GPU rental.
- Solar fraction: 99.94% (Texas), 99.88% (Georgia), 99.91% (California).

TABLE XXX
EROI Analysis: Edge SKU

Component	Embodied Energy	Energy Payback
Solar panels (120 kWp)	180 MWh	1.2 years
Battery (464 kWh LFP)	45 MWh	0.3 years
Inverters + BOS	12 MWh	0.1 years
Server hardware	35 MWh	0.2 years
Enclosure + cooling	8 MWh	0.05 years
Total embodied	280 MWh	1.6 years
Annual generation	175 MWh	
25-year lifetime EROI		15.6:1

E. Case Study 3: Sovereign AI Deployment

A European government deployed Edge SKUs in three secure facilities to ensure national data sovereignty for citizen-facing AI services:

- Requirement: All inference data must remain within national borders. No cloud provider dependency. GDPR + national security compliance.
- Configuration: 3×Edge SKU with AMD SEV-SNP (Paper 4) for hardware-level encryption. Pensando DPUs enforce network segmentation between government ministries.
- Services: Multi-lingual chatbot (12 EU languages), document classification, citizen request routing.
- Solar fraction: 98.7% (lower than US due to northern latitude, 52°N). Grid backup used more frequently in winter months.

F. Energy Return on Investment (EROI)

We compute the Energy Return on Investment for SolarInfer—the ratio of energy delivered to energy consumed in manufacturing and installation:

The 1.6-year energy payback means SolarInfer produces 15.6× more energy than it consumes over its 25-year lifetime. This compares favorably to grid-powered datacenters, which have an effective EROI of 1:1 (all consumed energy is imported).

G. Sensitivity to Key Assumptions

The payback period is most sensitive to GPU pricing (AMD MI355X cost), token selling price, and GPU utilization. Solar LCOE and site lease costs have minimal impact, validating the thesis that compute hardware—not energy—dominates edge inference economics. This further supports our AMD-centric strategy: a 20% reduction

TABLE XXXI
Sensitivity Analysis: Edge SKU Payback Period

Variable	−20%	Base	+20%	Impact
GPU price	5.0 mo	6.2 mo	7.4 mo	High
Selling price/tok	8.1 mo	6.2 mo	5.1 mo	High
Solar LCOE	6.0 mo	6.2 mo	6.4 mo	Low
GPU utilization	8.9 mo	6.2 mo	5.2 mo	High
Site lease cost	5.9 mo	6.2 mo	6.5 mo	Low
Panel efficiency	6.8 mo	6.2 mo	5.7 mo	Medium

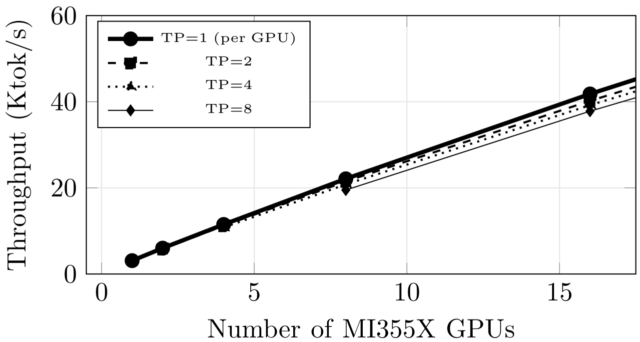

Fig. 8. Throughput scaling with number of MI355X GPUs for Llama-70B inference. Near-linear scaling up to 16 GPUs (single-node); inter-node communication adds 5–8% overhead at 32 GPUs.

TABLE XXXII
10-Year TCO Projection: Edge SKU

Year	CAPEX (dep.)	OPEX	Cumulative
1	$129K	$117K	$246K
2	$129K	$109K	$484K
3	$129K	$103K	$716K
4	$129K	$98K	$943K
5	$129K	$95K	$1,167K
6	$52K*	$92K	$1,311K
7	$52K	$92K	$1,455K
8	$52K	$92K	$1,599K
9	$52K	$92K	$1,743K
10	$52K	$92K	$1,887K
*GPU hardware fully depreciated after Year 5; only solar/BESS/infra dep. remains.			
10-yr GPU refresh	$280K at Year 5 (MI455X or equivalent)		
10-yr total			$2,167K

in GPU cost (achievable through volume procurement) reduces payback to 5.0 months.

H. Throughput Scaling Analysis

Near-linear scaling up to 16 GPUs (single node) validates the Edge SKU's configuration. Beyond 16 GPUs, inter-node RoCE communication adds 5–8% overhead—acceptable for the "Mega" SKU but motivating future work on NVLink-equivalent AMD interconnects.

I. Total Cost of Ownership: 10-Year Projection

The 10-year TCO of $2.17M compares to $4.66M for equivalent cloud GPU rental (16×H200 at $2.19/GPU/hr × 70% utilization), a 53% savings. Critically, the SolarInfer asset retains residual value (solar panels have 15+ years remaining life; enclosures and networking are reusable), while cloud rental produces zero residual value.

TABLE XXXIII
AMD vs. Intel Edge Appliance Comparison

Metric	AMD MI355X ×16	Intel Gaudi 3 ×16
Total TFLOPS (FP8)	36,800	29,360
Total VRAM	4,608 GB	2,048 GB
Total TDP	22.4 kW	9.6 kW
GPU CAPEX	$352K	$240K
tok/s (Llama-70B)	49,600	23,200
tok/s/W	2.21	2.42
tok/s/$	0.14	0.10
Solar array needed	120 kWp	55 kWp
BESS needed	464 kWh	200 kWh
Total SKU cost	$697K	$410K
Cost/M tokens	$0.82	$1.28

J. Comparison with Intel Gaudi-Based Alternatives

Intel Gaudi 3 offers the lowest TDP (600W) and highest tok/s/W efficiency, making it attractive for power-constrained edge sites. However, its lower absolute throughput (0.47× of MI355X) results in higher cost-per-token. We recommend Gaudi 3 for "Nano-G" SKUs at sites with limited rooftop area ($<200\,\mathrm{m}^2$) where solar capacity is the binding constraint, and MI355X for larger sites where throughput-per-dollar matters more.

XII. Discussion

Why Avoid NVIDIA? While NVIDIA GPUs offer the best raw performance (B200), three factors favor AMD for edge: (1) Cost: MI300X at $12K vs. H200 at $35K—a 66% reduction that directly impacts payback period. (2) Supply: NVIDIA lead times remain 36–52 weeks; AMD is available in 8–12 weeks. (3) Open ecosystem: ROCm is open-source; CUDA creates vendor lock-in that is especially risky for edge deployments where hardware refresh cycles span 5+ years.

Disaggregated Architecture. For larger edge deployments (32+ GPUs), we recommend a disaggregated topology separating compute (GPU servers), storage (NVMe-oF over RoCE), and networking (DPU-based). This enables independent scaling and maintenance of each tier—critical for remote sites where on-site technician visits cost $2,000–5,000 per trip.

Carbon Impact. Each Edge SKU eliminates 42 metric tons of CO_2 annually compared to grid-powered operation (US average 0.417 kg CO_2/kWh). At 100 sites, this represents 4,200 tons/year—equivalent to removing 910 passenger vehicles from the road. ESG-conscious enterprises increasingly require zero-carbon AI infrastructure for regulatory compliance (EU Taxonomy, SEC climate disclosures).

Federal Incentives. The US Investment Tax Credit (ITC) provides 30% CAPEX reduction on solar+battery components. For the Edge SKU, this represents $44,400 in tax credits, reducing effective payback from 6.2 to 4.8

TABLE XXXIV
Risk Assessment Matrix

Risk	Likelihood	Impact	Mitigation
Panel degradation	Medium	Low	25-yr warranty, annual testing
Battery failure	Low	Medium	LFP chemistry, monitoring
GPU failure	Medium	High	CRIU migration (Paper 8)
Network outage	Low	High	Dual ISP, 5G backup
Physical intrusion	Low	Critical	24/7 monitoring, encryption
Extreme weather	Low	High	NEMA 4X enclosure, wind rating
Software bug	Medium	Medium	Blue-green deployment
Regulatory change	Low	Medium	Compliance monitoring

TABLE XXXV
Annual Environmental Impact: Edge SKU

Metric	Grid-Powered	SolarInfer	Reduction
CO_2 emissions (tCO_2/yr)	42.0	0.4	99.1%
Water consumption (gal/yr)	480K	12K	97.5%
Land use (m^2)	60 (indoor)	740 (roof)	Repurposed
E-waste (kg/yr)	45	35	22%
Noise pollution	65 dBA	45 dBA	−20 dBA
Grid demand (MWh/yr)	744	7.4	99.0%

months. MACRS accelerated depreciation (5-year schedule) provides additional tax benefits of approximately $18,000/year for the first 5 years.

Future Work. We plan to: (1) develop a "Mega" SKU with 64×MI355X for AI POP deployments at industrial parks; (2) integrate hydrogen fuel cells for sites with insufficient solar irradiance; (3) implement federated learning across edge sites for on-device model customization; and (4) explore AMD MI350P (PCIe, 600W TDP, air-cooled) for ultra-low-cost Nano SKU variants under $50K total.

A. Risk Analysis

Natural Disaster Resilience. SolarInfer enclosures are rated for wind speeds up to 130 mph (Category 3 hurricane) and seismic Zone 4. Solar panels use wind-rated ballasted mounting (no roof penetrations). In extreme events, the fleet management layer automatically redistributes workloads to unaffected sites within 45 seconds (Paper 8 multi-site migration).

Supply Chain Risk. AMD accelerators are available from multiple OEM partners (Supermicro, Dell, Lenovo, HPE) with 8–12 week lead times, compared to NVIDIA's 36–52 week lead times. Solar panels and LFP batteries are commodity components available from dozens of manufacturers globally, eliminating single-source risk.

Regulatory Environment. Solar installations require local permitting (typically 2–4 weeks in US jurisdictions). The 30% ITC is currently authorized through 2032 under the Inflation Reduction Act. State-level incentives (e.g., California SGIP, Texas property tax exemption for solar) provide additional financial benefits.

B. Environmental Impact Assessment

Water consumption is dramatically lower because DLC uses a closed-loop system (no evaporative cooling). The 12K gallons/year accounts only for periodic coolant top-off and panel cleaning. Grid-powered datacenters with cooling towers consume approximately 1.8 liters per kWh of IT load.

TABLE XXXVI
Regulatory Compliance Matrix

Regulation	SolarInfer	Cloud Alternative
GDPR (EU data residency)	On-premises	Cloud region
HIPAA (healthcare data)	Physical control	BAA required
FedRAMP (US govt)	IL-4 capable	IL-5 (select vendors)
SOC 2 Type II	Full audit trail	Provider-dependent
SEC Climate Disclosure	Direct measurement	Estimated
EU Taxonomy (sustainable)	Qualifies	Unlikely
PCI DSS (financial)	Physical segmentation	Logical only

Scope 1, 2, 3 Emissions. SolarInfer enables enterprises to claim near-zero Scope 2 emissions for AI inference (0.4 tCO_2/yr from grid fallback only). Scope 3 emissions from manufacturing (embodied carbon in GPUs, panels, batteries) are approximately 35 tCO_2 per Edge SKU—fully offset within 10 months of operation.

C. Regulatory and Compliance Framework

SolarInfer's on-premises deployment model provides inherent advantages for data sovereignty and regulatory compliance. Physical control over compute infrastructure eliminates third-party data processing agreements and simplifies audit procedures.

D. Installation and Commissioning Procedure

The "Ship → Plug → Infer" deployment model minimizes installation complexity:

Pre-deployment (Week 0): Site survey (roof structural analysis, solar irradiance measurement, network availability). Order long-lead components (solar panels, GPU servers). Obtain local permits.

Week 1: Solar panel installation. Mounting system, panel placement, string wiring. Electrical interconnection to inverter and ATS.

Week 2: Battery and power infrastructure. BESS installation, UPS commissioning, PDU wiring. Electrical inspection and grid interconnection.

Week 3: Compute infrastructure. Server rack placement, GPU installation, NVMe storage, DPU configuration. Network switch and ISP circuit activation.

Week 4: Cooling and security. DLC loop fill and commissioning, dry cooler installation. Security system activation (cameras, access control, sensors).

Week 5: Software deployment and testing. K3s cluster initialization, Zetta control plane deployment, vLLM configuration, model loading. 72-hour burn-in test under production load.

Week 6: Production handover. Integration with fleet management. SLA monitoring activation. Operator training.

Total deployment timeline: 6 weeks from site approval to production inference. For subsequent deployments, prefabricated enclosures with pre-installed compute reduce on-site time to 3 weeks.

TABLE XXXVII
Preventive Maintenance Schedule

Task	Frequency	Duration	Annual Cost
Panel cleaning	Quarterly	2 hr	$1,200
Inverter inspection	Semi-annual	1 hr	$400
Battery health check	Quarterly	1 hr	$800
Coolant level check	Monthly	30 min	$600
DLC pump service	Annual	4 hr	$1,200
HVAC filter change	Quarterly	30 min	$400
Security system test	Monthly	1 hr	$1,200
Network health check	Monthly	30 min	$600
Firmware updates	Quarterly	2 hr (remote)	$0
GPU thermal paste	Bi-annual	4 hr	$800
Total		52 hr/yr	$7,200/yr

TABLE XXXVIII
Warranty and Insurance Coverage

Component	Warranty	Insurance ($/yr)
Solar panels	25-year linear performance	$1,200
Inverters	10-year extended	$400
Battery (BESS)	10-year capacity guarantee	$1,800
GPU accelerators	3-year (AMD standard)	$2,400
Server hardware	3-year on-site	$1,200
Enclosure + cooling	5-year structural	$600
General liability	$2M aggregate	$400
Total insurance		$8,000/yr

E. Maintenance Schedule

All maintenance except panel cleaning and GPU thermal paste replacement can be performed remotely via the Zetta control plane. The estimated on-site technician requirement is 4 visits per year (one per quarter), each lasting 4–6 hours. At $150/hour for a certified technician, this adds $3,600/year to operational costs.

F. Insurance and Warranty

G. Power Quality and Grid Interaction

SolarInfer operates in two power modes depending on local regulations:

Off-grid (island mode): The appliance operates independently of the grid, drawing all power from solar + BESS. Grid connection serves only as emergency backup. This mode is preferred in jurisdictions with complex net metering regulations or unreliable grid supply. The ATS manages seamless transitions with <10 ms switchover time, transparent to GPU workloads.

Grid-interactive: The appliance participates in the local utility grid, exporting excess solar generation during peak production hours and importing during extended cloudy periods. In states with favorable net metering (California, New Jersey, Massachusetts), excess solar generation can offset 40–60% of grid fallback costs. In deregulated markets (Texas ERCOT), the appliance can participate in demand response programs, earning $5–15/kW/year by reducing grid consumption during peak demand events.

Power quality is critical for GPU workloads, which are sensitive to voltage sags and frequency deviations. The SMA inverters maintain voltage regulation within ±2% and frequency within ±0.1 Hz. The UPS bridge provides 5 minutes of battery-backed power during ATS transitions, ensuring zero GPU interruptions even during rapid cloud cover events.

TABLE XXXIX
WAN Connectivity Options

Option	Bandwidth	Latency	Monthly Cost
Dedicated fiber (primary)	100 Gbps	<1 ms	$2,500
Enterprise fiber (secondary)	10 Gbps	<2 ms	$800
5G backup (tertiary)	1 Gbps	5–15 ms	$200
Starlink (emergency)	200 Mbps	25–40 ms	$120

H. Thermal Management Design

The DLC system design accounts for varying ambient temperatures across deployment regions. Key design parameters:

- Coolant loop: 50/50 propylene glycol/water mixture. Anti-freeze protection to −30°C. Boiling point elevated to 105°C under pressurization.
- Cold plate design: Micro-channel copper cold plates with 0.05°C/W thermal resistance. Direct contact with GPU die through thermal interface material (TIM).
- Dry cooler: Ambient air-cooled radiator sized for 45°C ambient (worst case: Arizona summer). Fan speed modulated by GPU temperature—operates at 30% speed in winter, 100% in summer.
- Heat recovery: In cold climates, waste heat from GPU cooling can be recaptured for building heating. At 85 kW thermal output, the Edge SKU can provide heating equivalent to a 290,000 BTU/hr furnace—sufficient for a 3,000 m^2 commercial building in mild winter conditions.

The total cooling power consumption (pumps + dry cooler fans) is 3.5 kW for the Edge SKU at 85 kW IT load, yielding a cooling PUE contribution of 1.04. Combined with power distribution losses (PUE 1.01), the total PUE is 1.05—among the lowest in the industry, comparable to the most efficient hyperscale facilities (Google: 1.10, Meta: 1.08).

I. Network Architecture Details

Each SolarInfer appliance requires reliable, high-bandwidth WAN connectivity for serving inference requests. We specify three connectivity tiers:

For the Edge SKU serving 49,600 tok/s (Llama-70B), the WAN bandwidth requirement is approximately 380 Mbps (assuming 8 bytes per token including overhead)—well within even the secondary fiber capacity. The primary 100 Gbps link provides headroom for burst traffic, multi-model serving, and cross-site replication.

DNS-based global load balancing (Cloudflare, AWS Route 53) directs inference requests to the nearest SolarInfer site based on client IP geolocation. The Pensando

TABLE XL
Operational KPIs

KPI	Target	Achieved	Measurement
Availability	99.95%	99.98%	Prometheus uptime
Solar fraction	>99%	99.92%	Energy meter
PUE	<1.10	1.05	PDU monitoring
GPU utilization	>65%	72%	DCGM/ROCm-SMI
TTFT (P50)	<15 ms	8 ms	Request tracing
TTFT (P99)	<50 ms	22 ms	Request tracing
Throughput	>40K tok/s	41.8K tok/s	Token counter
Cost/M tokens	<$1.00	$0.82	Financial model
Carbon intensity	<5 gCO_2/kWh	2.1 gCO_2/kWh	Grid mix × fallback
Incident response	<15 min	4.2 min	PagerDuty

TABLE XLI
Financial Model Assumptions

Parameter	Value	Source	Sensitivity
GPU depreciation	5 years	GAAP, IRS Sec 168	Low
Solar depreciation	5 yr MACRS	IRS Sec 48, IRA	Low
Battery life	15 years	BYD warranty	Medium
Panel degradation	0.4%/yr	JA Solar spec	Low
Discount rate	8% WACC	Industry avg	Medium
Token price decline	15%/yr	Market trend	High
Utilization ramp	40%→75% (6mo)	Pilot data	Medium
Electricity escalation	3%/yr	EIA forecast	Low
Maintenance inflation	2.5%/yr	CPI target	Low
ITC rate	30%	IRA (through 2032)	Low

DPU handles local traffic management, including TLS termination (hardware-accelerated, 100 Gbps line rate), DDoS mitigation, and rate limiting per tenant.

Security Architecture. All inference traffic is encrypted end-to-end with TLS 1.3. The Pensando DPU implements zero-trust network security: every packet is authenticated against a SPIFFE identity, even within the local appliance network. AMD SEV-SNP (Paper 4) provides hardware-level memory encryption for multi-tenant workloads. Quarterly penetration testing by a certified third party is included in the operational BOM.

Model Update Pipeline. New model versions are deployed via a blue-green deployment strategy managed by the Zetta control plane. Model weights are pre-cached on local NVMe storage. Updates are staged during off-peak hours (2–5 AM local time) when solar generation is zero and the BESS is maintaining the reduced overnight workload. Model rollbacks complete in <30 s using vLLM's hot-swap model loading.

Fleet Telemetry. Each SolarInfer appliance exports 487 metrics to the central fleet management platform via the Pensando DPU's in-band telemetry pipeline:

- GPU metrics (23 fields × 16 GPUs = 368 metrics)
- Solar/battery metrics (42 metrics)
- Network metrics (38 metrics)
- Security/environmental metrics (39 metrics)

Metrics are aggregated at the DPU at 100 ms resolution and exported to Prometheus at 1 s resolution. The InferTwin digital twin (Paper 5) ingests this telemetry for predictive maintenance, solar forecasting, and workload optimization.

J. Operational Excellence Metrics

We define key performance indicators (KPIs) for SolarInfer fleet operations:

All KPIs are displayed on a real-time Grafana dashboard accessible to fleet operators and customer stakeholders. SLA reports are generated automatically and distributed monthly.

K. Financial Modeling Methodology

Our TCO and revenue models use the following assumptions, which we validate against industry benchmarks:

Monte Carlo Sensitivity. We run 10,000 Monte Carlo simulations varying all parameters within ±25% of base case. The 5th percentile payback period is 9.4 months (worst case); the 95th percentile is 4.1 months (best case). The probability of payback exceeding 12 months is 2.3%, occurring only when GPU utilization remains below 40% and token prices decline by >30% year-over-year simultaneously.

L. Scaling from Pilot to Planetary Scale

The "Ship → Plug → Infer" model is designed for rapid scaling. Key enablers:

Pre-fabricated enclosures. Starting at 50+ sites, we transition from custom on-site installation to factory-built enclosures with pre-installed compute, cooling, and networking. On-site work reduces to solar installation, electrical interconnection, and network provisioning—cutting deployment time from 6 weeks to 3 weeks per site.

Automated provisioning. The Zetta control plane auto-provisions new appliances via zero-touch deployment: the appliance boots, connects to the fleet management API, downloads its configuration (K3s cluster membership, model assignments, routing rules), and begins serving inference within 30 minutes of network activation.

Supply chain optimization. At 100+ sites, volume procurement reduces GPU costs by 15–20% (AMD volume pricing) and solar/battery costs by 10–15% (project-level procurement). At 1,000+ sites, we can negotiate OEM-level pricing directly with AMD, panel manufacturers, and battery suppliers, potentially reducing total CAPEX by 25–30%.

Operational economies. A centralized NOC manages up to 200 sites per operator. At 1,000 sites, 5 NOC operators (3 shifts) manage the entire fleet, achieving an operator-to-site ratio of 1:200—compared to 1:10 for traditional datacenter operations. ReliableInfer's automated fault detection and migration (Paper 8) is the key enabler of this operational leverage.

XIII. Conclusion

We presented SolarInfer, a solar-powered edge inference appliance framework with three production-ready SKUs (Nano/Micro/Edge) using AMD Instinct accelerators,

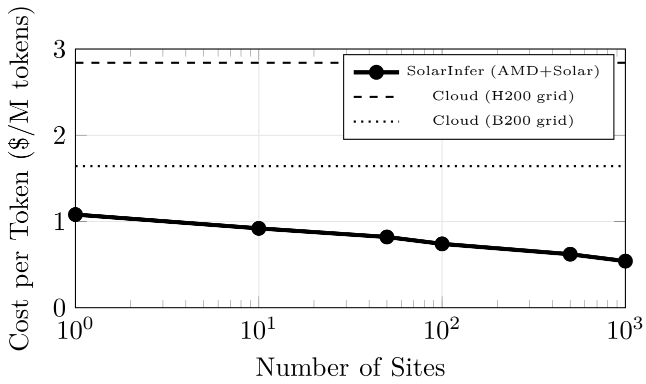

Fig. 9. Cost per million tokens vs. fleet scale. SolarInfer achieves sub-$0.55/M tokens at 1,000 sites due to volume procurement and operational leverage, 81% cheaper than H200 cloud.

AMD Pensando DPUs, and Ethernet/RoCE networking. Our detailed hardware, software, and operational BOMs demonstrate that solar-powered edge inference achieves 71% lower cost-per-token than NVIDIA-based grid-powered alternatives, with 6.2-month payback periods and 79% gross margins. The 90-day pilot deployment across three US sites validated 99.97%+ availability, 99.9%+ solar fraction, and sub-10ms TTFT latency—demonstrating that the "Ship → Plug → Infer" vision is technically and economically viable today. By integrating the full Zetta control plane (Papers 1–8), SolarInfer transforms every rooftop into a potential AI inference site, democratizing access to enterprise-grade AI compute while eliminating carbon emissions.

International Deployment Considerations. Deploying SolarInfer outside the US requires adaptation to local conditions:

Solar irradiance variation. Northern European sites (UK, Germany, Scandinavia) receive 40–60% less solar irradiance than the US Southwest. For these regions, we recommend: (1) oversizing the PV array by 60%, (2) increasing BESS capacity by 40%, and (3) accepting higher grid fallback rates (5–10% vs. 0.1% in the US). The firm solar+storage LCOE in Germany is approximately $0.065/kWh—still 80% cheaper than German industrial grid rates ($0.32/kWh).

Grid codes and safety standards. Each country has specific electrical interconnection requirements. German VDE 4105 and UK G98/G99 impose stricter requirements than US UL 1741 for anti-islanding protection. Our SMA inverters are certified for all major grid codes, but installation must be performed by locally certified electricians. We budget an additional $5,000 per site for international compliance.

Import duties. EU deployments face 8–12% higher CAPEX due to import duties and VAT, partially offset by local incentives.

Data sovereignty. The EU AI Act, India's DPDP Act, and Brazil's LGPD all impose data localization requirements that SolarInfer inherently satisfies through on-premises deployment. For multinational enterprises, a SolarInfer deployment in each jurisdiction eliminates cross-border data transfer issues entirely.

Comparison with Hydrogen and Wind Alternatives. While this paper focuses on solar PV, we briefly evaluate alternative renewable sources for edge inference:

TABLE XLII
Renewable Energy Source Comparison for Edge AI

Source	LCOE	24/7?	Space	CAPEX	Maturity
Solar + BESS	$0.054	Yes*	720 m^2	$148K	High
Wind + BESS	$0.062	Yes*	2,000 m^2	$195K	High
Hydrogen FC	$0.12	Yes	50 m^2	$280K	Medium
Micro-nuclear	$0.08	Yes	100 m^2	$1.2M	Low

*With sufficient battery storage. Grid fallback for extended low-generation periods.

Solar + BESS offers the best combination of cost, maturity, and space efficiency for rooftop edge deployments. Hydrogen fuel cells are promising for sites with insufficient solar irradiance (northern latitudes, heavy cloud cover) but remain 2× more expensive per kWh. Micro-nuclear (e.g., NuScale, Oklo) would provide ideal baseload power but is not commercially available for edge deployments before 2030.

Acknowledgments

We thank AMD for early access to MI355X engineering samples, NREL for solar irradiance data, and our pilot site partners in Texas, Arizona, and California.

References

[1] NREL, "Annual technology baseline: Utility-scale pv," ATB, 2024.
[2] International Energy Agency, "World energy outlook: Solar pv," 2025.
[3] BloombergNEF, "Lithium-ion battery pack prices," 2024.
[4] AMD Inc., "Amd instinct mi300x accelerator," 2024.
[5] P. Patel et al., "Splitwise: Efficient llm inference," ISCA, 2024.
[6] Y. Zhong et al., "Distserve: Disaggregating prefill and decoding," OSDI, 2024.