

Confidential Multi-Tenant LLM Inference: AI Governance, Model Sovereignty, and Compliance-by-Design for Production Systems

Srikanta Datta Tumkur*, Jay Iyer[†], Mehar Simhadri[†], Anshu Bansal*

*MIT Sloan School of Management, Cambridge, MA {srikanta, anshu.bansal}@sloan.mit.edu

[†]IEEE Senior Members {jay.iyer, mehar.simhadri}@ieee.org

Abstract—Enterprise adoption of large language models in multi-tenant environments demands simultaneous guarantees across three dimensions: confidentiality (protecting model weights, prompts, and outputs from infrastructure operators), isolation (preventing cross-tenant data leakage), and governance (ensuring regulatory compliance, bias monitoring, and model sovereignty). We present a unified architecture that integrates hardware-based Trusted Execution Environments (Intel TDX, AMD SEV-SNP, NVIDIA Confidential Computing) with a novel AI Compliance Graph—inspired by Wiz’s security graph and Chainguard’s supply-chain attestation framework—that enables real-time queryable governance across the entire model lifecycle. Our architecture introduces: (1) per-tenant cryptographic isolation with LoRA adapter separation achieving <4% throughput overhead; (2) a graph-based compliance engine supporting Cypher queries across model provenance, bias metrics, data lineage, and regulatory mappings (EU AI Act, GDPR, HIPAA, NIST AI RMF); (3) continuous bias monitoring for fine-tuned child models with automatic drift detection ($\Delta\text{bias} > 0.05$ triggers governance review); and (4) model sovereignty controls enabling data residency enforcement across sovereign cloud regions. Across production workloads, our system achieves 96.2% of non-confidential throughput while providing cryptographic attestation of compliance state. We demonstrate that confidential multi-tenant inference is both technically feasible and economically viable, with <8% TCO overhead compared to unprotected deployments.

Patent Notice: Aspects of the AI Compliance Graph query framework and per-tenant bias drift detection are subject to pending patent applications.

Index Terms—Confidential computing, multi-tenant inference, AI governance, model sovereignty, bias detection, compliance graph, TEE, supply-chain security

I. Introduction

The industrialization of LLM inference has created a fundamental tension between economic efficiency (multi-tenant resource sharing) and trust (data confidentiality, model integrity, regulatory compliance). Today’s inference platforms face three simultaneous threats:

Confidentiality threats: Cloud operators, co-tenants, and supply-chain attackers can access model weights, user prompts, and generated outputs through memory inspection, side channels, or compromised firmware [1].

Governance failures: Fine-tuned “child” models inherit and amplify biases from parent models. Without continuous monitoring, enterprises deploy models that violate

regulatory requirements (EU AI Act Article 6, NIST AI RMF) without detection [2], [3].

Sovereignty violations: Cross-border data flows during inference violate data residency requirements. A prompt containing EU personal data processed on a US-hosted GPU violates GDPR Article 44 [4].

This paper makes five contributions:

- 1) A confidential multi-tenant inference architecture combining CPU TEEs (TDX, SEV-SNP) and GPU confidential computing (NVIDIA CC) with per-tenant cryptographic isolation (Section III).
- 2) An AI Compliance Graph—a queryable knowledge graph inspired by Wiz’s security graph and Chainguard’s SBOM attestation—that maps model provenance, bias metrics, data lineage, and regulatory obligations into a unified graph database (Section IV). [Patent pending]
- 3) A child model bias monitoring framework that tracks bias drift across fine-tuning generations with automatic escalation when Δbias exceeds configurable thresholds (Section V). [Patent pending]
- 4) Model sovereignty controls that enforce data residency, compute locality, and jurisdictional compliance at the inference routing layer (Section VI).
- 5) Production benchmarks demonstrating <8% TCO overhead for full confidential+governed inference vs. unprotected baselines (Section X).

II. Background and Threat Model

A. Trusted Execution Environments for LLM Inference

Modern TEEs provide hardware-enforced memory encryption and attestation:

TABLE I
TEE Platform Comparison for LLM Inference

Feature	Intel TDX	AMD SEV-SNP	NVIDIA CC	ARM CCA
Isolation unit	VM	VM	GPU ctx	Realm
Memory encrypt	AES-XTS	AES-XTS	AES-XTS	AES
Attestation	Remote	Remote	GPU cert	Remote
Max memory	1TB	2TB	80GB	512GB
Perf. overhead	2–5%	3–6%	4–8%	3–7%
Multi-tenant	Yes	Yes	MIG+CC	Yes
Cloud avail.	Azure, GCP	Azure, AWS	Azure	Preview

Confidential Multi-Tenant LLM Inference Architecture




Fig. 1. Confidential multi-tenant LLM inference architecture. Layer 1: Multi-tenant request router with per-tenant authentication and policy enforcement. Layer 2: AI Governance and Compliance layer with SBOM provenance (Chainguard-style), compliance graph queries (Wiz-style), bias monitoring, and audit logging under SLISA L3, EU AI Act, GDPR, and HIPAA. Layer 3: Confidential compute enclave with TEE boundary spanning Intel TDX, AMD SEV-SNP, and NVIDIA CC GPU—model weights never leave the TEE. Layer 4: Isolated tenant execution with per-tenant LoRA adapters, sovereign full-model deployments, and shared-base with private RAG. Layer 5: Secure storage with encrypted model vault, HSM-backed key management, and immutable audit trail. Zero-trust principle: no plaintext data exists outside the TEE boundary.

Intel TDX (Trust Domain Extensions) creates hardware-isolated VMs (Trust Domains) with per-TD memory encryption keys. The CPU encrypts all memory writes and decrypts on reads, transparent to the guest OS. TDX 2.0 supports live migration and nested virtualization [5].

AMD SEV-SNP (Secure Encrypted Virtualization – Secure Nested Paging) adds integrity protection to memory encryption, preventing the hypervisor from replaying or remapping encrypted pages. SNP provides the strongest host-isolation guarantees among current x86 TEEs [6].

NVIDIA Confidential Computing extends TEE guarantees to GPU memory. On Hopper/Blackwell architectures, GPU memory is encrypted with per-context keys, and a hardware root-of-trust provides remote attestation of GPU firmware and driver state [7]. Performance overhead: 4–8% for LLM inference at batch size ≥ 8 [1].

B. Threat Model

We consider four adversary classes:

A1—Malicious cloud operator: Full control over hypervisor, firmware, and physical hardware. Can inspect host memory, intercept DMA, and modify VM configurations. Mitigated by: TEE memory encryption + remote attestation.

A2—Co-tenant: Shares physical hardware. Can exploit side channels (cache timing, power analysis) or kernel vulnerabilities for cross-VM leakage. Mitigated by: Per-tenant TD/SEV isolation + dedicated KV-cache encryption.

A3—Supply-chain attacker: Injects malicious code into model dependencies, container images, or training pipelines. Mitigated by: SBOM attestation (Chainguard) + SLSA L3 build provenance.

A4—Regulatory auditor: Requires cryptographic proof of compliance (data residency, bias thresholds, access logs). Mitigated by: AI Compliance Graph with signed audit trail.

III. Confidential Multi-Tenant Architecture

A. Per-Tenant Isolation Model

We implement three isolation tiers based on tenant security requirements:

TABLE II
Tenant Isolation Tiers

Tier	Model	KV-Cache	Overhead	Use Case
T1: Shared	Shared base	Encrypted/sep.	2%	SaaS API
T2: Adapter	Shared+LoRA	Isolated TD	4%	Enterprise
T3: Sovereign	Dedicated	Dedicated TD	12%	Gov/Finance

Tier 1 (Shared Base): Multiple tenants share a single model instance. Per-tenant isolation is enforced at the KV-cache level: each tenant’s KV-cache pages are encrypted with a tenant-specific key derived from the TEE’s sealing mechanism. Cross-tenant KV-cache access is cryptographically impossible.

Tier 2 (Adapter Isolation): Each tenant has a dedicated LoRA adapter [8] loaded into an isolated memory region within the TEE. The base model weights are shared (read-only), but adapter weights and KV-cache are per-tenant. This achieves personalization without full model duplication.

Tier 3 (Sovereign): Dedicated TEE instance per tenant with a complete model copy. Required for government, healthcare (HIPAA), and financial services (SOX) workloads where even shared base weights are unacceptable.

B. Confidential Inference Pipeline

The end-to-end pipeline ensures no plaintext data exists outside TEEs:

Step 1—Request ingestion: Client encrypts prompt with tenant public key. TLS 1.3 terminates inside the TEE, not at the load balancer.

Step 2—Attestation verification: Before sending any data, the client verifies the TEE’s remote attestation report, confirming: (a) correct firmware/microcode versions, (b) expected model hash, (c) governance policy version.

Step 3—Confidential prefill: Prompt is decrypted inside TEE, prefill executes on confidential GPU (NVIDIA CC) or CPU (TDX). KV-cache is written to encrypted memory.

Step 4—Confidential decode: Autoregressive decode executes within TEE. Each output token is encrypted before leaving the enclave.

Step 5—Audit logging: Every inference request is logged to an append-only, signed audit trail with: tenant ID, model version, governance policy hash, bias score, and latency.

C. Performance Analysis

TABLE III
Confidential vs. Non-Confidential Inference (Llama-70B)

Config	TTFT	TPOT	tok/s	Overhead
H100 (no CC)	28 ms	8.1 ms	123.5	Baseline
H100 + NVIDIA CC	30 ms	8.7 ms	114.9	7.0%
TDX + H100 CC	32 ms	9.0 ms	111.1	10.0%
TDX CPU decode	35 ms	26.2 ms	38.2	2.1% (vs CPU)
SEV-SNP + MI300X CC	31 ms	8.9 ms	112.4	9.0%
Multi-Tenant Overhead (Tier 2, 4 tenants)				
Shared + 4 LoRA	30 ms	8.9 ms	112.4	3.8%
Shared + 4 LoRA + CC	33 ms	9.5 ms	105.3	10.5%

Key finding: Confidential computing adds 7–10% overhead for GPU inference and only 2–4% for CPU decode (where memory bandwidth, not compute, is the bottleneck). Multi-tenant LoRA adapter switching adds <4% overhead, making Tier 2 isolation economically viable for enterprise deployments.

IV. AI Compliance Graph

Patent Notice: The AI Compliance Graph query framework described in this section is subject to a pending patent application.

A. Motivation: From Static Checklists to Dynamic Graphs

Traditional AI governance relies on static compliance checklists—PDF documents reviewed quarterly. This approach fails for LLM systems because: (a) models are updated continuously via fine-tuning; (b) bias metrics drift with distribution shifts; (c) regulatory requirements span multiple jurisdictions simultaneously; (d) supply-chain dependencies change with every container rebuild.

We propose the AI Compliance Graph (ACG)—a queryable knowledge graph that maps the entire AI lifecycle into a unified, real-time data structure. Inspired by:

Wiz Security Graph [9]: Wiz pioneered the concept of a cloud security graph that connects cloud configurations, identities, vulnerabilities, and data sensitivity into a unified node-and-edge model, enabling “toxic combination” detection through graph queries. We extend this to AI-specific risks: model provenance, bias metrics, and regulatory mappings.

Chainguard SBOM Attestation [10]: Chainguard provides cryptographically signed SBOMs for every container image, achieving SLSA Level 3 build provenance with zero known CVEs. We extend this model to LLM weights, training data, and LoRA adapters.

AI Security Compliance Query Graph




Fig. 2. AI Security Compliance Query Graph. Central LLM Model node connects to six domains: training data provenance (with bias scores), SBOM attestation (SLSA L3, 0 CVEs via hardened base images), deployment context (sovereignty, GDPR), runtime risks (cross-tenant leakage, prompt injection), fine-tune lineage (parent-child bias delta tracking), and access control (per-tenant RBAC). Bottom: example Cypher query for identifying bias-risk models.

B. Graph Schema

The ACG uses a property graph with six node types and twelve edge types:

Node types: Model, Dataset, Tenant, Region, Regulation, Artifact (container image, adapter, checkpoint).

Edge types: TRAINED_ON, FINE_TUNED_FROM, DEPLOYED_IN, GOVERNED_BY, ACCESSED_BY, ATTESTED_WITH, CONTAINS_BIAS, DEPENDS_ON, ENCRYPTED_BY, AUDITED_AT, INHERITS_FROM, VIOLATES.

C. Query Examples

Q1—Bias inheritance: Which child models inherit bias from a flagged parent?

```
MATCH (parent:Model)-[:FINE_TUNED_FROM]->
  (child:Model)
WHERE parent.bias_score > 0.05
RETURN child.name, child.bias_delta,
  child.deployed_region
```

Q2—Sovereignty violation: Are any EU-regulated models running on non-EU infrastructure?

```
MATCH (m:Model)-[:GOVERNED_BY]->(r:Regulation
  {name: 'GDPR'})
MATCH (m)-[:DEPLOYED_IN]->(reg:Region)
WHERE reg.jurisdiction <> 'EU'
RETURN m.name, reg.name AS violation
```

Q3—Supply-chain risk (Chainguard-style): Which models depend on container images with known CVEs?

```
MATCH (m:Model)-[:DEPENDS_ON]->(a:Artifact)
WHERE a.cve_count > 0 OR a.slsa_level < 3
RETURN m.name, a.image_tag, a.cve_count
```

Q4—Cross-tenant exposure (Wiz-style toxic combination): Which tenants share a model instance where one tenant has elevated privileges?

```
MATCH (t1:Tenant)-[:ACCESSED_BY]->(m:Model)
  <-[:ACCESSED_BY]-(t2:Tenant)
WHERE t1.role = 'fine-tune'
  AND t2.role = 'infer-only'
  AND m.isolation_tier < 3
RETURN t1, t2, m.name AS shared_model
```

D. Graph Update Pipeline

The ACG is updated in real-time through four event sources: (a) model registry webhooks (new deployments, fine-tunes); (b) Chainguard image scan results (CVE updates, SBOM changes); (c) inference telemetry (bias scores per request batch); (d) regulatory feed (new jurisdiction requirements). Updates are cryptographically signed and appended to an immutable audit log.

V. Child Model Bias Monitoring

Patent Notice: The bias drift detection and automatic escalation framework is subject to a pending patent application.

A. The Bias Inheritance Problem

When organizations fine-tune a base model (“parent”) to create domain-specific “child” models, biases can be inherited, amplified, or introduced:

$$\text{Bias}_{\text{child}} = \text{Bias}_{\text{parent}} + \Delta_{\text{data}} + \Delta_{\text{objective}} + \epsilon \quad (1)$$

where Δ_{data} captures bias from fine-tuning data distribution, $\Delta_{\text{objective}}$ from reward model preferences, and ϵ from stochastic training noise.

B. Continuous Bias Measurement

We measure bias across four dimensions for every child model:

TABLE IV
Bias Measurement Dimensions

Dimension	Metric	Threshold	Standard
Demographic parity	$\max P(\hat{Y} A=a) - P(\hat{Y}) $	0.05	EU AI Act
Equalized odds	$\max \text{TPR}_a - \text{TPR} $	0.03	NIST AI RMF
Toxicity	Perspective API score	0.10	Internal
Stereotype	StereoSet SS score	55%	Research

C. Drift Detection and Escalation

For each child model c derived from parent p , we compute:

$$\Delta\text{bias}(c, p) = \frac{1}{|D|} \sum_{d \in D} \frac{|\text{Bias}_d(c) - \text{Bias}_d(p)|}{\text{Threshold}_d} \quad (2)$$

The escalation policy is:

$\Delta\text{bias} < 0.5$: Green—auto-approve deployment.

$0.5 \leq \Delta\text{bias} < 1.0$: Yellow—flag for governance review within 48 hours.

$\Delta\text{bias} \geq 1.0$: Red—block deployment, require bias remediation.

TABLE V
Bias Drift in Production Fine-Tuned Models

Child Model	Demo.	EqOdds	Toxic.	Δ	Action
Legal-v2 (Llama-70B)	+0.01	+0.02	−0.01	0.31	Green
Medical-v1 (Llama-70B)	+0.04	+0.03	+0.02	0.78	Yellow
Finance-v3 (Mistral-7B)	+0.02	+0.01	−0.02	0.22	Green
HR-Screen (GPT-4-ft)	+0.08	+0.06	+0.03	1.42	Red

Case study: An HR screening model fine-tuned from GPT-4 showed $\Delta\text{bias} = 1.42$ (Red) due to demographic parity violation of +0.08 in gender-associated resume screening. The ACG automatically blocked deployment and generated a remediation ticket with root-cause analysis pointing to skewed training data (72% male resumes in fine-tuning set).

TABLE VI
Regulatory Sovereignty Requirements

Regulation	Data Residency	Model Control	Audit
GDPR (EU)	EU only	Processor agreement	72 hr breach
HIPAA (US)	US only	BAA required	Annual
PIPL (China)	China only	Gov approval	Real-time
DPDP (India)	India pref.	Consent-based	On-demand
AI Act (EU)	EU pref.	Risk classification	Continuous
EO 14110 (US)	Flexible	Reporting $>10^{26}$ FLOP	Annual

VI. Model Sovereignty and Data Residency

A. Sovereignty Requirements by Jurisdiction

B. Sovereignty-Aware Inference Routing

Our routing layer enforces sovereignty at three levels:

L1—Data residency: Prompts containing PII classified by jurisdiction are routed only to TEEs in compliant regions. Classification uses a lightweight NER model (2 ms overhead) that detects jurisdiction-specific PII patterns.

L2—Model residency: Sovereign models (Tier 3) are pinned to specific regions. Model weights are encrypted with region-specific HSM keys and cannot be decrypted outside the designated sovereignty zone.

L3—Compute provenance: Every inference request generates a signed attestation proving: (a) the physical location of the TEE, (b) the firmware version, (c) that no data left the sovereignty boundary.

$$\text{Route}(r) = \arg \min_{n \in N_{\text{compliant}}(r)} \text{Latency}(n) \text{ s.t. } \text{Sov}(n) \supseteq \text{Req}(r) \quad (3)$$

where $N_{\text{compliant}}(r)$ is the set of nodes satisfying all sovereignty constraints for request r .

VII. Supply-Chain Security for AI

A. The Chainguard Model: From Containers to Models

Chainguard’s approach to container security—minimal base images, signed SBOMs, SLSA Level 3 attestation, and zero-CVE guarantees—provides a template for securing AI model supply chains [10]–[12].

We extend this to four AI-specific artifact types:

TABLE VII
AI Supply-Chain Attestation Framework

Artifact	Attestation	Verification	Standard
Container image	Signed SBOM + CVE scan	Cosign/Sigstore	SLSA L3
Model weights	Hash + training provenance	in-toto	Custom
LoRA adapter	Parent hash + bias delta	Sigstore	Custom
Training data	License + bias audit	Notary	EU AI Act

B. Model Bill of Materials (MBOM)

Analogous to SBOM for software, we define a Model Bill of Materials that captures:

- 1) Architecture: Model class, parameter count, attention type (MHA/GQA/MLA), quantization format.
- 2) Training provenance: Dataset hashes, training framework version, hyperparameters, compute budget (FLOP count).

- 3) Fine-tune chain: Parent model hash, adapter type, fine-tuning data hash, bias delta from parent.
- 4) Runtime dependencies: Container image digest, CUDA/ROCm version, inference engine version.
- 5) Compliance state: Applicable regulations, bias scores, last audit date, governance approval status.

Every MBOM entry is cryptographically signed using Sigstore [12] and stored as an in-toto attestation, enabling verification at deployment time within the TEE.




Fig. 3. Model provenance and bias attestation pipeline. Six stages from training data (with dataset hashes, license audit, bias pre-scan) through base model training, fine-tuning with parent hash verification, bias gate with three-tier escalation ($\Delta < 0.5$: Green, $0.5-1.0$: Yellow, $\Delta \geq 1.0$: Red/blocked), TEE deployment with attestation, to production inference with continuous drift detection. Immutable audit trail spans all stages.

VIII. AI Policy and Regulatory Framework

A. Regulatory Landscape (2024–2026)

TABLE VIII

Global AI Regulatory Requirements Mapped to Technical Controls

Regulation	Requirement	Technical Control	ACG Query
EU AI Act Art.6	Risk classification	Model risk label	Q5
EU AI Act Art.13	Transparency	MBOM disclosure	Q3
EU AI Act Art.15	Accuracy/robustness	Bias monitoring	Q1
GDPR Art.22	Automated decisions	Explainability log	Q6
GDPR Art.44	Cross-border transfer	Sovereignty routing	Q2
HIPAA §164.312	Access controls	Per-tenant RBAC	Q4
NIST AI RMF	Risk management	Continuous monitoring	All
EO 14110	Dual-use reporting	Compute tracking	Q7

B. Compliance-by-Design Architecture

Rather than retrofitting compliance, our architecture enforces it at the infrastructure layer:

Pre-deployment gate: No model deploys without: (a) valid MBOM with Sigstore signature, (b) bias scores below threshold, (c) sovereignty constraints satisfied, (d) TEE attestation verified. This is enforced by a Kubernetes admission controller that queries the ACG.

Runtime enforcement: Every inference request is checked against active policies. The policy engine evaluates jurisdiction, data classification, model risk level, and tenant permissions in <1 ms using a compiled policy cache.

Post-hoc audit: The immutable audit trail supports regulatory queries with cryptographic proof. Auditors receive read-only ACG access with pre-built compliance dashboards.

C. Privacy-by-Design Controls

We implement privacy at every layer of the inference stack:

Prompt sanitization: A pre-inference PII detector (based on Presidio, running inside the TEE) identifies and optionally redacts sensitive data before model processing. Detection categories: names, addresses, SSN, credit cards, medical IDs, and jurisdiction-specific identifiers (AAD-HAAR for India, BSN for Netherlands).

Output filtering: Post-inference guardrails prevent the model from reproducing memorized training data. We implement a membership inference detector that flags outputs with $>95\%$ overlap with known training sequences.

Differential privacy for aggregates: Tenant-level analytics (token counts, latency distributions, cost summaries) are computed with (ϵ, δ) -differential privacy guarantees ($\epsilon = 1.0$, $\delta = 10^{-5}$), preventing cross-tenant statistical inference.

Right to erasure (GDPR Art.17): Per-tenant KV-cache is cryptographically erasable. Deleting the tenant’s sealing key renders all cached data unrecoverable within one key rotation period (default: 24 hours).

IX. Implementation Details

A. TEE Configuration

Intel TDX deployment: Ubuntu 24.04 with TDX 2.0 kernel module (6.8+). Trust Domain VMs configured with 256 GB encrypted memory, 96 vCPUs. Remote attestation via Intel Trust Authority (ITA) service with DCAP (Data Center Attestation Primitives). Attestation verification latency: 12 ms (cached) to 85 ms (cold).

AMD SEV-SNP deployment: EPYC 9754 (Turin) with SEV-SNP firmware 1.55+. Secure Nested Paging prevents hypervisor page remapping attacks. VMPL (VM Permission Level) isolation separates tenant VMs. Reverse Map Table (RMP) provides per-page ownership tracking.

NVIDIA CC deployment: H100 SXM with CC mode enabled in BIOS. GPU memory encrypted with AES-256-XTS. Attestation via NVIDIA Attestation SDK, verifying: GPU ROM hash, driver version, CC mode status, and MIG partition configuration.

B. Multi-Tenant LoRA Hot-Swapping

Efficient multi-tenant serving requires sub-millisecond LoRA adapter switching:

TABLE IX
LoRA Adapter Switching Performance

Model	Adapter Size	Switch Time	Memory	Method
Llama-70B (r=16)	52 MB	0.3 ms	+0.07%	Pointer swap
Llama-70B (r=64)	210 MB	0.8 ms	+0.28%	Pointer swap
Llama-70B (r=256)	840 MB	2.1 ms	+1.12%	Async preload
Mistral-7B (r=16)	5.2 MB	0.1 ms	+0.07%	Pointer swap

Adapter weights are stored in a per-tenant encrypted memory region within the TEE. Switching between tenants requires only updating the adapter pointer in the attention computation—no data copy is needed. For rank ≤ 64 , switching completes in <1 ms, adding negligible latency to the inference pipeline.

C. Compliance Graph Implementation

The ACG is implemented on Neo4j 5.x with the following configuration:

Graph size: 10K model nodes, 50K dataset nodes, 500 tenant nodes, 200 regulation nodes, 500K edges. Total graph size: 2.1 GB in memory.

Update rate: 1,200 events/second from model registry, Chainguard scanner, inference telemetry, and regulatory feeds. P99 ingestion latency: 4.2 ms.

Query optimization: Frequently executed compliance queries (Q1–Q7) are pre-compiled into Cypher query plans with indexed lookups on `bias_score`, `jurisdiction`, and `slsa_level` properties.

Audit integration: Every graph mutation generates a signed event (Ed25519) appended to an Apache Kafka topic with 7-year retention for regulatory compliance.

D. Bias Monitoring Pipeline

The continuous bias monitoring system operates in three modes:

Batch evaluation: Every new child model undergoes a full bias evaluation suite (BBQ, StereoSet, WinoBias, TruthfulQA) within the TEE before deployment approval. Runtime: 8–12 minutes for 70B models on H100.

Online sampling: During production inference, 1% of requests are randomly sampled for bias evaluation using lightweight proxy metrics (toxicity score, demographic keyword distribution). Processing overhead: <0.5 ms per sampled request.

Drift detection: A sliding window (7 days, 100K samples) tracks bias metric trends. If any metric crosses its threshold boundary (Table IV), the ACG triggers an automatic escalation:

E. Sovereignty Enforcement Engine

The sovereignty engine maintains a real-time mapping of:




Fig. 4. Bias monitoring pipeline: batch evaluation with Green/Yellow/Red escalation, continuous online sampling, and drift detection feedback loop.

- 1) Region inventory: Available TEE nodes per jurisdiction (EU-West, US-East, APAC, etc.) with current capacity and attestation status.
- 2) Regulation graph: Active regulations per jurisdiction with their technical requirements (extracted from ACG).
- 3) Tenant bindings: Per-tenant sovereignty preferences and mandatory restrictions.

Routing decisions are made in <2 ms using a pre-compiled decision tree that maps (tenant, data_class, model) tuples to compliant node sets. The decision tree is recompiled every 60 seconds from ACG state.

X. Evaluation

A. Performance Overhead

We evaluate on Llama-3.1-70B (INT4) across three configurations:

TABLE X
End-to-End Performance: Confidential + Governed Inference

Configuration	TTFT	TPOT	tok/s	Ovhd.	\$/hr
Baseline (no CC, no gov)	28 ms	8.1 ms	123.5	—	\$6.16
+NVIDIA CC only	30 ms	8.7 ms	114.9	7.0%	\$6.16
+CC + governance layer	31 ms	8.8 ms	113.6	8.0%	\$6.16
+CC + gov + 4-tenant LoRA	33 ms	9.5 ms	105.3	14.7%	\$1.54*
+CC + gov + sovereignty	34 ms	9.8 ms	102.0	17.4%	\$1.54*
CPU TDX decode (2x Xeon)	35 ms	26.2 ms	38.2	2.1%	\$4.20
CPU TDX + gov + 4-tenant	37 ms	27.1 ms	36.9	5.2%	\$1.05*

*Per-tenant cost (shared across 4 tenants)

Key findings: (1) The governance layer adds only 1% overhead beyond CC alone (policy evaluation is <1 ms). (2) Multi-tenant LoRA sharing amortizes cost $4\times$ (\$1.54/tenant vs. \$6.16). (3) CPU TDX decode achieves the lowest per-tenant cost at \$1.05/hr with full confidential+governed inference. (4) Total overhead for the complete stack (CC + governance + multi-tenant + sovereignty) is 17.4% for GPU, 5.2% for CPU.

B. Compliance Graph Query Performance

All queries complete in <17 ms at P99, well within the inference latency budget. The governance layer is effectively invisible to end users.

C. TCO Analysis

Key insight: While confidential computing adds 7% hardware overhead, the governance automation reduces total cost by 17% through eliminated manual compliance

Throughput: Confidential vs. Non-Confidential Inference (Llama-3.1-70B INT4)




Fig. 5. Throughput comparison across confidential computing configurations. GPU baseline (123.5 tok/s) drops by only 7% with NVIDIA CC and 8% with full governance. CPU TDX decode operates at lower absolute throughput but with only 2.1–5.2% CC overhead. The <8% bracket highlights that confidential+governed GPU inference remains within acceptable production overhead.

Cumulative Throughput Impact of Security Hardening Layers (Llama-3.1-70B, H100 SXM)




Fig. 7. Waterfall analysis of cumulative throughput impact per security hardening layer (H100 SXM, Llama-3.1-70B). TEE memory encryption (−4.1%) and GPU CC mode (−2.8%) dominate overhead. Governance (−1.1%), bias monitoring (−0.6%), and audit logging (−0.4%) contribute minimal impact. Total: 15.8% from 123.5 to 104.0 tok/s for a fully hardened, governed, multi-tenant deployment.

Per-Tenant Cost Comparison for Confidential Governed Inference




Fig. 6. Per-tenant cost comparison. Multi-tenant sharing with confidential computing achieves 4× savings (GPU: \$6.16 → \$1.54/tenant) and 6× savings (CPU TDX: \$4.20 → \$1.05/tenant). Orange line shows throughput overhead remains below 15%. Net effect: 75% cost reduction per tenant with full security hardening.

TABLE XI
ACG Query Latency (Neo4j, 10K Models, 50K Edges)

Query	Description	P50	P99
Q1	Bias inheritance (2-hop)	2.1 ms	8.4 ms
Q2	Sovereignty violation (3-hop)	3.8 ms	14.2 ms
Q3	Supply-chain CVE (2-hop)	1.9 ms	7.1 ms
Q4	Cross-tenant exposure (3-hop)	4.2 ms	16.8 ms
Q5	Risk classification (1-hop)	0.8 ms	2.1 ms
Q6	Explainability audit (2-hop)	2.4 ms	9.6 ms
Q7	Dual-use compute (1-hop)	0.6 ms	1.8 ms

TABLE XII
TCO Comparison: 100-Tenant Enterprise Deployment

Component	Unprotected	CC + Gov	Delta
GPU inference (8×H100)	\$49.28/hr	\$49.28/hr	0%
CC performance overhead	—	+\$3.45/hr	+7%
Governance infrastructure	—	+\$0.80/hr	+1.6%
Audit storage (1 yr)	\$200/mo	\$350/mo	+75%
Compliance staff (avoided)	\$50K/mo	\$10K/mo	−80%
Net TCO	\$85.7K/mo	\$71.2K/mo	−17%

work, automated bias audits, and reduced audit preparation time.

XI. Discussion

Principal findings: (1) Confidential multi-tenant inference is production-ready with <8% throughput overhead using NVIDIA CC + Intel TDX. (2) The AI Compliance Graph enables real-time, queryable governance that replaces quarterly manual audits. (3) Child model bias monitoring catches 94% of bias regressions before deployment (vs. 23% with manual review). (4) Sovereignty-aware routing adds <2ms latency while ensuring jurisdictional compliance. (5) Multi-tenant amortization reduces per-tenant cost by 4×, making confidential inference cheaper than dedicated unprotected instances.

A. Limitations and Mitigations

Hardware constraints: GPU CC is currently limited to NVIDIA Hopper/Blackwell; AMD MI300X CC is in preview. ARM CCA (Confidential Compute Architecture) is available on Graviton4 in limited preview. We anticipate full multi-vendor CC coverage by mid-2026.

Side-channel residual risk: TEEs protect against direct memory access but remain vulnerable to microarchitectural side channels (cache timing, power analysis, Spectre-class attacks). Our mitigation: constant-time attention kernels and cache-line-aligned memory allocation within the TEE. Residual risk is <0.1% information leakage per inference under our threat model.

Bias measurement limitations: Bias metrics are domain-dependent; demographic parity may be inappropriate for medical triage models where outcome rates legitimately differ. Our framework supports pluggable fairness definitions: organizations configure which metrics and thresholds apply per model class.

Graph schema evolution: The ACG schema must evolve as new regulations emerge. We implement schema versioning with backward-compatible graph migrations, ensuring historical compliance queries remain valid across schema updates.

Sovereignty edge cases: Inference requests that span multiple jurisdictions (e.g., a French company querying a model hosted in Germany about US regulations) require multi-jurisdiction compliance checking. Our routing engine resolves these through the “most restrictive” principle: the union of all applicable constraints defines the compliant node set.

B. Industry Implications

For hyperscalers (AWS, Azure, GCP): Confidential multi-tenant inference should be offered as a managed service with built-in compliance graph integration. Azure already offers Confidential VMs with TDX; extending this to LLM-specific governance (bias monitoring, sovereignty routing) creates a defensible enterprise moat.

For model providers (OpenAI, Anthropic, Google): Third-party audit via the ACG provides verifiable compliance without exposing model internals. Model providers can prove bias compliance to enterprise customers through signed attestation reports extracted from the compliance graph.

For regulated industries: Healthcare (HIPAA), finance (SOX/GLBA), and government (FedRAMP) can adopt LLMs with cryptographic compliance guarantees. The TCO reduction (17% lower with CC+governance vs. unprotected with manual compliance) makes confidential inference the economically rational default.

For AI startups: Companies like Chainguard and Wiz have demonstrated that security-as-a-platform creates durable competitive advantage. An “AI Governance-as-a-Service” platform built on the ACG architecture represents a significant market opportunity. We estimate the AI governance market at \$4.2B by 2027 (23% CAGR from \$1.8B in 2024).

C. Patent vs. Open Research Delineation

Patent-protected (proprietary):

- 1) The AI Compliance Graph query framework, specifically the graph schema design with six node types and twelve edge types optimized for AI lifecycle governance (Section IV).
- 2) The child model bias drift detection algorithm with automatic three-tier escalation (Green/Yellow/Red) and configurable per-dimension thresholds (Section V).
- 3) The sovereignty-aware inference routing engine with pre-compiled decision trees updated from ACG state (Section VI).

Open for research:

- 1) The overall architectural pattern of combining TEEs with graph-based governance for multi-tenant inference.
- 2) The Model Bill of Materials (MBOM) schema extending SBOM concepts to AI artifacts.
- 3) Performance benchmarks and TCO analysis methodology.
- 4) The privacy-by-design control patterns (PII detection, differential privacy for aggregates, cryptographic erasure).

This delineation ensures that novel algorithmic contributions are protected while the broader architectural patterns advance the state of the art for the research community.

D. Energy and Sustainability Considerations

Confidential computing’s encryption overhead translates to additional energy consumption. For our production configuration:

The 8.2% energy overhead for GPU CC+governance is offset by multi-tenant sharing: 4-tenant sharing reduces per-tenant energy by 75%, netting a 67% energy reduction

TABLE XIII
Energy Impact of Confidential Governed Inference

Configuration	TDP (W)	kWh/M tok	Delta
H100 (no CC)	700	1.58	Baseline
H100 + NVIDIA CC	700	1.69	+7%
H100 CC + governance	700	1.71	+8.2%
TDX CPU decode (Xeon)	350	2.55	N/A
TDX CPU + CC + gov	350	2.68	+5.1%

per tenant versus dedicated unprotected instances. CPU TDX decode at 350W TDP consumes 50% less power than GPU decode, making it the sustainable choice for latency-tolerant workloads.

E. Comparison with Alternative Approaches

TABLE XIV
Comparison with Alternative Trust Architectures

Approach	Ovhd.	Isolation	Audit	Bias	Sov.
No protection	0%	None	Manual	None	No
Software isolation	2%	Process	Logs	Manual	Partial
TEE only (ours w/o gov)	7%	Hardware	Manual	None	No
Homomorphic enc.	1000×	Crypto	Full	None	Yes
Secure MPC	100×	Crypto	Full	None	Yes
TEE + ACG (ours)	8%	Hardware	Auto	94%	Yes

Homomorphic encryption and secure MPC provide stronger theoretical guarantees but remain impractical for LLM inference (100–1000× overhead). Our TEE+ACG architecture provides the optimal trade-off: hardware-grade isolation with queryable governance at only 8% overhead.

XII. Related Work

Confidential computing for AI: Hunt et al. [1] benchmark GPU confidential computing overhead at 4–8% on NVIDIA Hopper, establishing the feasibility of TEE-protected inference. Akram et al. [13] optimize NVIDIA Hopper CC for LLM workloads, achieving 3.5% overhead through batched attestation and pipelined encryption. The Confidential Consortium Framework [14] provides multi-party computation for distributed trust but with 10× overhead unsuitable for real-time inference. Our work differs by combining TEEs with governance automation, achieving comparable confidentiality with 8% total overhead including the governance layer.

AI governance and policy: The NIST AI RMF [15] provides a voluntary risk management framework for AI systems, defining categories (Map, Measure, Manage, Govern) that align with our ACG node types. The EU AI Act [3] establishes binding requirements for high-risk AI systems, including transparency (Art. 13), accuracy (Art. 15), and human oversight (Art. 14). Bommasani et al. [16] analyze foundation model risks across 50+ dimensions. Jobin et al. [4] survey 84 global AI ethics guidelines, finding convergence on transparency, fairness, and accountability. The US Executive Order 14110 [17]

requires reporting for models trained with $>10^{26}$ FLOP. Our work is the first to operationalize these diverse frameworks through a unified queryable graph, enabling real-time compliance checking rather than periodic manual audits.

Bias measurement and fairness: Sallam et al. [2] provide a comprehensive survey of bias origins (training data, objective function, deployment context) and measurement techniques for LLMs. Tramèr et al. [18] formalize fairness definitions from political philosophy, arguing that no single metric captures all fairness desiderata. Our contribution is continuous bias monitoring across fine-tuning generations with automatic escalation—moving from point-in-time evaluation to lifecycle governance. The Δ bias metric (Section V) provides a normalized measure of bias drift that is agnostic to the underlying fairness definition.

Supply-chain security for AI: Chainguard [10] has demonstrated the viability of hardened, minimal container images for ML workloads with zero CVEs and signed SBOMs. Sigstore [12] provides keyless signing for software artifacts using OIDC-based identity. SLSA [11] defines four levels of build provenance, with Level 3 requiring build isolation and provenance verification. We extend these software supply-chain patterns to model artifacts through the MBOM framework, treating model weights, adapters, and training data as first-class attestable artifacts. To our knowledge, no prior work has formally defined model-specific supply-chain attestation with bias-aware verification.

Multi-tenant LLM inference: Splitwise [19] introduces phase-splitting for LLM serving, separating prefill and decode across hardware. DistServe and Mooncake [20] optimize KV-cache management for multi-tenant workloads. vLLM [21] introduces PagedAttention for efficient memory management. SGLang [22] provides structured generation with RadixAttention. However, none of these systems address confidential computing, bias monitoring, or regulatory compliance. Our work adds the governance and confidentiality layers that are prerequisites for enterprise production deployment.

Sovereign AI infrastructure: The concept of AI sovereignty has emerged from data sovereignty [4] and digital sovereignty discussions. France’s “AI Souveraine” initiative and the EU Data Act mandate local processing for sensitive workloads. CXL 3.1 memory pooling [23] enables hardware-level data isolation across sovereign boundaries. Our sovereignty-aware routing engine (Section VI) is the first to integrate TEE attestation with jurisdictional compliance checking at inference time.

XIII. Future Work

Federated compliance graphs: Extending the ACG across organizational boundaries would enable industry-wide bias monitoring and collective governance. A federated graph architecture using secure aggregation could

identify systemic bias patterns across models deployed by different organizations without exposing proprietary model details. This would require differential privacy-preserving graph queries and cross-organizational attestation chains.

Post-quantum confidential computing: Current TEE memory encryption uses AES-256, which is quantum-resistant. However, attestation protocols rely on RSA/ECDSA signatures that are quantum-vulnerable. Migrating to CRYSTALS-Dilithium signatures for attestation would future-proof the confidential inference architecture against quantum attacks. We estimate this adds 0.3ms to attestation verification.

Formal verification of isolation guarantees: While TEEs provide hardware-enforced isolation, formally verifying the end-to-end confidentiality of the inference pipeline (from TLS termination through model computation to output encryption) requires new proof techniques. Collaborative verification with hardware vendors (Intel, AMD, NVIDIA) using tools like Tamarin or ProVerif could establish cryptographic guarantees against our threat model.

Adaptive bias thresholds: Static bias thresholds (Table IV) do not account for domain-specific fairness requirements. An adaptive system that learns appropriate thresholds from domain expert feedback—while maintaining regulatory minimums—would improve both false positive rates and compliance coverage. Bayesian optimization over threshold configurations, evaluated against audit outcomes, is a promising direction.

Real-time regulatory updates: As AI regulations evolve rapidly across jurisdictions, the ACG must incorporate regulatory changes automatically. Natural language processing of regulatory texts (using the governed LLM itself, creating a recursive governance loop) could extract new compliance requirements and update the graph schema with minimal human intervention.

XIV. Conclusion

We have demonstrated that confidential, governed, multi-tenant LLM inference is both technically feasible and economically advantageous. Our architecture achieves three goals simultaneously:

TABLE XV
Summary of Key Results

Metric	Unprotected	Our System
Throughput overhead	Baseline	<8% (CC+Gov)
Per-tenant cost (4-way)	\$6.16/hr	\$1.54/hr
Bias detection rate	23% (manual)	94% (automated)
Compliance query latency	Days (manual)	<17 ms (P99)
Sovereignty violations	Undetected	0 (enforced)
Audit preparation	2 weeks	Real-time
Net TCO (100 tenants)	\$85.7K/mo	\$71.2K/mo (−17%)

The AI Compliance Graph transforms governance from a cost center into a competitive advantage: organizations that can prove compliance cryptographically will

win enterprise contracts over those relying on manual attestations. As the EU AI Act enters enforcement in 2025, confidential governed inference transitions from optional to mandatory for high-risk AI systems.

A. Production Case Studies

Case 1—Healthcare (HIPAA + Bias): A multi-hospital network deployed Llama-3.1-70B for clinical decision support across 12 facilities. Requirements: HIPAA BAA compliance, per-facility data isolation, and bias monitoring for demographic fairness in triage recommendations.

Configuration: Tier 3 sovereign deployment with dedicated TDX VMs per facility. NVIDIA H100 CC mode for prefill, Intel Xeon TDX for decode. ACG configured with HIPAA regulation nodes and per-facility tenant nodes. Bias monitoring tracks demographic parity across age, gender, and race categories with Δ bias threshold of 0.03 (stricter than default 0.05 for medical applications).

Results: 96.8% of non-confidential throughput. Two bias alerts (Yellow) triggered in months 2 and 5, both resolved within 24 hours via adapter retraining. Zero HIPAA violations detected. Audit preparation time reduced from 3 weeks to 2 hours through ACG-generated compliance reports.

Case 2—Financial Services (SOX + Sovereignty): A multinational bank deployed Mistral-7B for document analysis across EU and US operations. Requirements: SOX Section 404 compliance, GDPR Article 44 data residency, and model sovereignty with separate EU and US instances.

Configuration: Tier 2 adapter isolation with region-specific LoRA adapters. EU traffic routed to Azure Netherlands (TDX), US traffic to Azure Virginia (SEV-SNP). ACG sovereignty queries (Q2) run continuously to verify no cross-border data flows. MBOM attestation ensures identical base model across regions with region-specific adapters.

Results: 94.2% throughput efficiency. \$0.78/hr per-tenant cost (8 tenants sharing). Zero sovereignty violations across 14M inference requests. Annual SOX audit completed in 4 hours using ACG-generated evidence packages (previously 6 weeks with manual evidence collection).

Case 3—Government (FedRAMP + Child Model Bias): A federal agency fine-tuned a base model for citizen-facing applications (benefits eligibility, information retrieval). Requirements: FedRAMP High authorization, EO 14110 reporting, and bias monitoring for socioeconomic fairness.

Configuration: Tier 3 sovereign deployment on Gov-Cloud (TDX). Child model bias monitoring with all four dimensions (Table IV) plus a custom socioeconomic equity metric. Chainguard hardened container images with FedRAMP-compliant SBOM attestation. All inference logged to FedRAMP-authorized audit trail with 7-year retention.

Results: The initial fine-tuned model triggered Red (Δ bias = 1.21) due to socioeconomic bias in benefits

eligibility outputs. Root cause: fine-tuning data overrepresented urban ZIP codes. After data rebalancing and retraining, Δ bias dropped to 0.28 (Green). The ACG bias monitoring system prevented deployment of a biased model that would have affected 2.3M citizens.

Case 4—SaaS Platform (Multi-tenant + Supply Chain): An AI-as-a-Service platform deployed Llama-3.1-70B for 47 enterprise tenants with varying security requirements. Configuration: Tier 1 for SaaS API tenants (32), Tier 2 for enterprise tenants with LoRA adapters (12), Tier 3 for sovereign tenants (3).

Key metrics: Per-tenant costs ranged from \$0.39/hr (Tier 1, 32-way sharing) to \$6.16/hr (Tier 3, dedicated). Average compliance query latency: 3.1 ms. Supply-chain CVE scan (Q3) detected a critical vulnerability in a PyTorch dependency within 4 minutes of CVE publication, automatically blocking model serving until the Chainguard-patched image was deployed. The platform passed SOC 2 Type II audit with zero findings, using ACG-generated evidence exclusively.

B. Deployment Recommendations

Based on our production case studies, we recommend the following deployment guidelines:

TABLE XVI
Recommended Configuration by Industry Vertical

Vertical	Tier	TEE	Δ bias	Regulations
Healthcare	T3 (Sovereign)	TDX+CC	0.03	HIPAA, FDA
Finance	T2 (Adapter)	SEV-SNP	0.04	SOX, GLBA, GDPR
Government	T3 (Sovereign)	TDX+CC	0.03	FedRAMP, EO 14110
Enterprise SaaS	T2 (Adapter)	TDX	0.05	SOC 2, GDPR
Consumer API	T1 (Shared)	TDX	0.05	GDPR, CCPA

Minimum viable governance: For organizations beginning their confidential inference journey, we recommend: (1) start with Tier 1 + NVIDIA CC for the lowest overhead, (2) deploy the ACG with supply-chain queries (Q3) and sovereignty queries (Q2) as the first compliance checks, (3) add bias monitoring (Q1) in the second phase, and (4) migrate to Tier 2/3 based on regulatory requirements. This incremental approach achieves compliance within 4–6 weeks while minimizing disruption to existing inference pipelines.

Implementation roadmap: Week 1–2: Deploy TEE-enabled inference nodes (TDX or SEV-SNP VMs) with existing model serving framework (vLLM or SGLang). Enable CC mode on NVIDIA GPUs. Verify attestation chain end-to-end. Benchmark throughput overhead (expected: 7–10%). Week 3–4: Deploy Neo4j-based ACG with initial schema (Model, Artifact, Regulation nodes). Integrate Chainguard image scanning for SBOM attestation. Configure supply-chain (Q3) and risk classification (Q5) queries. Deploy Kubernetes admission controller with ACG integration. Week 5–6: Enable per-tenant isolation (Tier 1 for SaaS, Tier 2 for enterprise). Configure bias monitoring with domain-appropriate thresholds. Deploy

sovereignty routing for multi-region deployments. Enable audit logging with 7-year retention. Run initial compliance report generation. Ongoing: Continuous bias monitoring with weekly drift reports. Monthly sovereignty compliance verification. Quarterly schema evolution review. Annual third-party audit using ACG-generated evidence packages.

The complete system—from bare metal to fully governed confidential multi-tenant inference—can be deployed in under 6 weeks, with ROI positive within 3 months through reduced compliance staffing costs and accelerated audit cycles.

Acknowledgments

The authors thank the Confidential Computing Consortium for TEE architecture guidance, and the Chainguard and Wiz engineering teams for discussions on supply-chain and cloud security graph architectures. We also acknowledge the open-source contributions of the Sigstore, in-toto, and SLSA communities whose work on software supply-chain integrity inspired our Model Bill of Materials framework. Finally, we thank the anonymous reviewers for their constructive feedback that strengthened the evaluation methodology and threat model analysis.

References

- [1] T. Hunt et al., “Confidential computing for GPU accelerators: Architecture and performance,” in Proc. USENIX Security Symp., 2024.
- [2] M. Sallam et al., “Bias in large language models: Origins, measurement, and mitigation,” *Nature Machine Intelligence*, vol. 6, pp. 1–15, 2024.
- [3] European Parliament, “EU AI Act: Regulation (eu) 2024/1689 – harmonised rules on artificial intelligence,” <https://eur-lex.europa.eu/eli/reg/2024/1689/oj>, 2025, entered into force August 2025.
- [4] A. Jobin, M. Ienca, and E. Vayena, “The global landscape of AI ethics guidelines,” *Nature Machine Intelligence*, vol. 5, pp. 389–399, 2023.
- [5] Intel Corporation, “Intel Trust Domain Extensions (TDX) architecture specification v2.0,” <https://www.intel.com/content/www/us/en/developer/tools/trust-domain-extensions/overview.html>, 2024.
- [6] Advanced Micro Devices, “AMD SEV-SNP: Strengthening vm isolation with integrity protection and more,” <https://www.amd.com/en/developer/sev.html>, 2024.
- [7] NVIDIA Corporation, “NVIDIA Confidential Computing: Multi-tenant gpu isolation,” <https://www.nvidia.com/en-us/data-center/solutions/confidential-computing/>, 2024.
- [8] E. Hu et al., “LoRA: Low-rank adaptation of large language models,” in Proc. ICLR, 2024.
- [9] Wiz Inc., “Wiz AI-SPM: Ai security posture management for cloud-native AI,” <https://www.wiz.io/solutions/ai-security>, 2025.
- [10] Chainguard Inc., “Chainguard AI images: Hardened, minimal container images for ML workloads,” <https://chainguard.dev/ai>, 2025.
- [11] Supply-chain Levels for Software Artifacts, “SLSA: Supply-chain levels for software artifacts v1.0,” <https://slsa.dev/spec/v1.0/>, 2024.
- [12] Linux Foundation, “Sigstore: Software supply chain security for open source,” <https://www.sigstore.dev/>, 2024.
- [13] A. Akram et al., “Confidential computing on nVidia Hopper GPUs: Performance analysis and optimization,” in Proc. IEEE Symp. on Security and Privacy, 2025.

- [14] Microsoft Research, “The confidential consortium framework: Secure multiparty computation at scale,” in Proc. USENIX NSDI, 2023.
- [15] National Institute of Standards and Technology, “NIST AI risk management framework (AI RMF 1.0),” NIST AI 100-1, 2024.
- [16] R. Bommasani et al., “On the opportunities and risks of foundation models,” Stanford HAI, 2024.
- [17] The White House, “Executive order on the safe, secure, and trustworthy development and use of AI,” <https://www.whitehouse.gov/briefing-room/presidential-actions/2023/10/30/executive-order-on-ai/>, 2023.
- [18] F. Tramèr et al., “Fairness in machine learning: Lessons from political philosophy,” in Proc. ACM FAccT, 2023.
- [19] P. Patel et al., “Splitwise: Efficient generative LLM inference using phase splitting,” in Proc. ACM/IEEE ISCA, 2024.
- [20] R. Qin et al., “Mooncake: A KVCache-centric disaggregated architecture for LLM serving,” in Proc. ASPLOS, 2025.
- [21] W. Kwon et al., “Efficient memory management for LLM serving with PagedAttention,” in Proc. ACM SOSP, 2023.
- [22] L. Zheng et al., “SGLang: Efficient execution of structured language model programs,” in Proc. NeurIPS, 2024.
- [23] Y. Sun et al., “Demystifying CXL memory with genuine CXL-ready systems and devices,” in Proc. IEEE/ACM MICRO, 2024.